

Array Navigation Learning Model

10/14/2021

In the following code, we investigate species level differences in learning rates of four species of primates. The task the primates are learning is to visit all locations within local arrays shaped like a double trapezoid, a pentagon, and a Z. We model fit a model of the total distance an individual animal travels before visiting all locations relative to the number of trials that individual has undergone, the species of the individual, and the individual's own initial performance and learning rate. Because we expect learning curves to follow a power-law function, we log-log transform the distance traveled and trial number. Prior to running this analysis, the script “Traplining_Data_Processing_Script.R” was run to prepare the data. For each array, we also used simulations of agents moving between locations in identical arrays according to an iterative reinforcement algorithm, and compared the results of simulations in each array to the corresponding primate data. The results of this analysis are described in the manuscript “Do primates really trapline? Mild recursion and evidence for variable cognitive mechanisms in five species”.

First, we load the data and conduct log transformations

```
knitr::opts_chunk$set(echo = FALSE)
setwd("C:/Users/avining/Documents/Manuscripts/Do Primates Trapline")

#enable easy changes in text size, for tidy LaTeX in pdf
def.chunk.hook <- knitr::knit_hooks$get("chunk")
knitr:::knit_hooks$set(chunk = function(x, options) {
  x <- def.chunk.hook(x, options)
  ifelse(options$size != "normalsize",
        paste0("\n \\", options$size, "\n\n", x, "\n\n \\normalsize"),
        x)
})

library(tidyverse)
library(rstan)
library(bayesplot)
library(RColorBrewer)

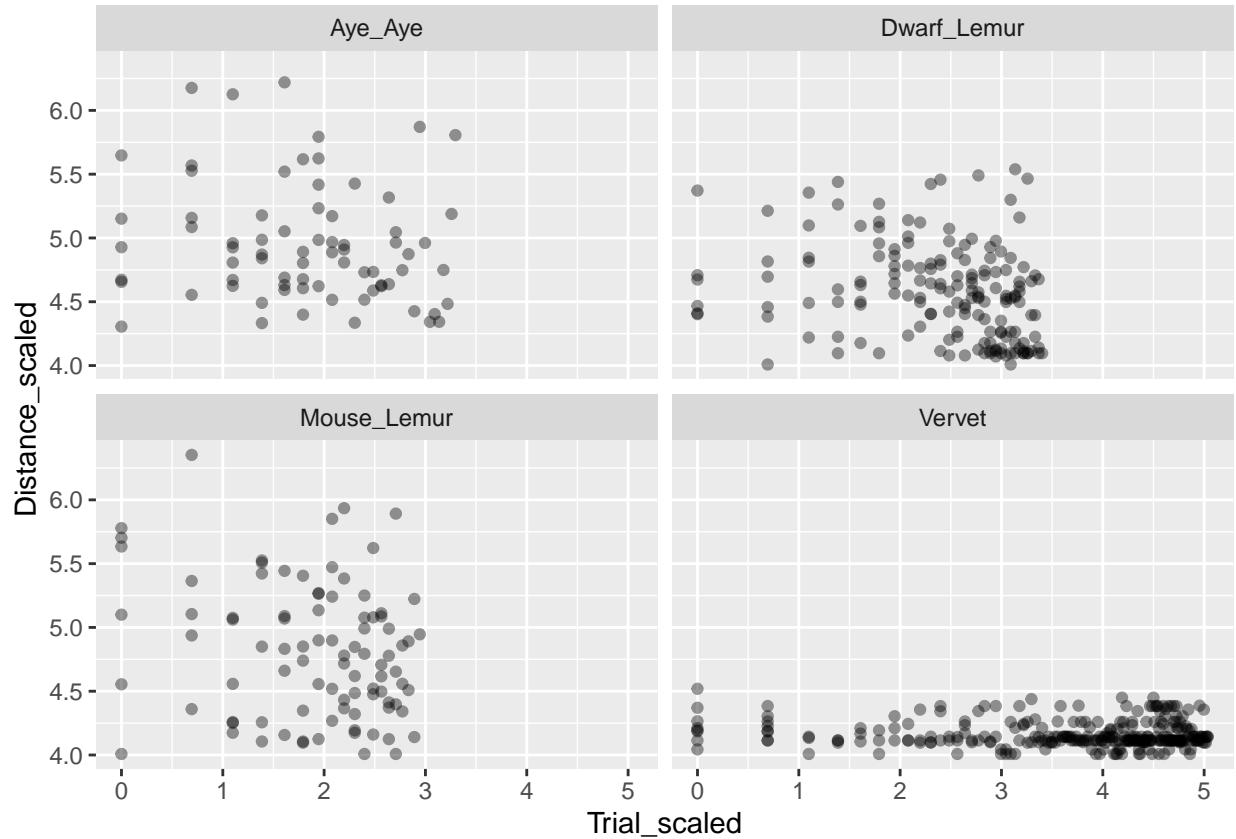
data_cleaned <- read_csv("./Results/data_cleaned.csv", col_types = "dccccdddddcc")

#log/log transform data to model a power function using linear fit
data_cleaned$Distance_scaled <- log(data_cleaned$Distance)
data_cleaned$Trial_scaled <- log(data_cleaned$Trial)
data_cleaned$Source <- gsub("Learning", "Reinforcement", data_cleaned$Source)
```

Then, we analyze the results from each array separately, starting with the Double Trapezoid. Trials were conducted in the this way with aye ayes, dwarf lemurs, mouse lemurs, and vervets.

Double Trapezoid (DT) Analysis

We begin by visualizing the data and perparing for analysis in STAN



In the above plot we see the log-log scaled data through which we will attempt a linear fit (analogous to a power function fit through the raw data). They do appear to show decreasing trends in all species but the vervet, which exhibits shorter distances traveled on initial trials. To investigate these trends, we fit the data to our model described in the manuscript “Do primates really trapline? Mild recursion and evidence for variable cognitive mechanisms in five species”.

Model Diagnostics

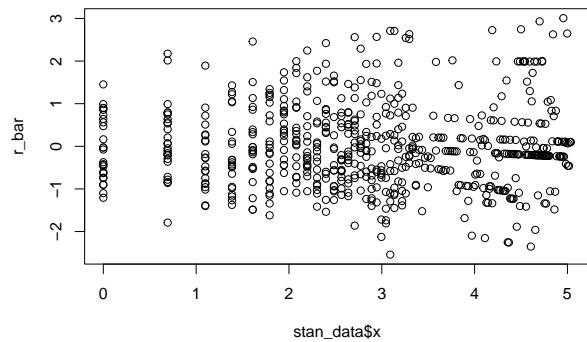
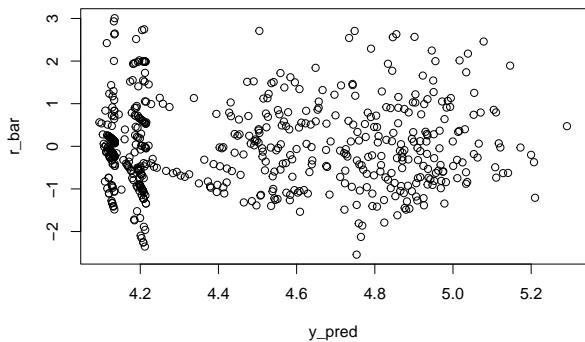
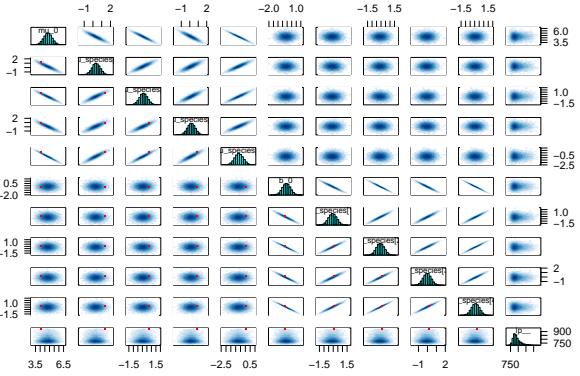
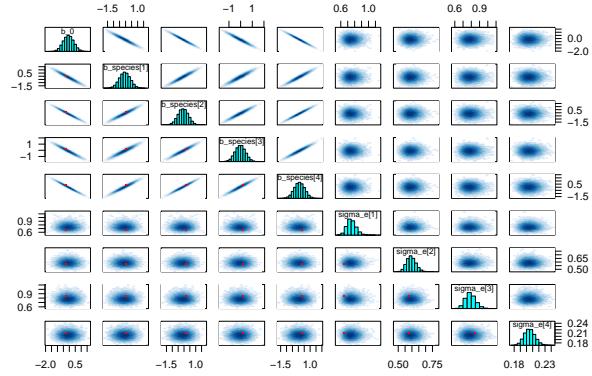
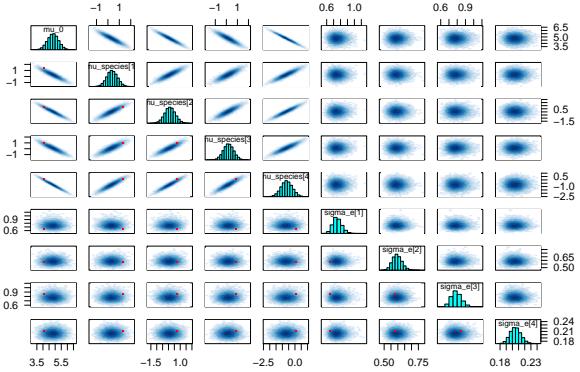
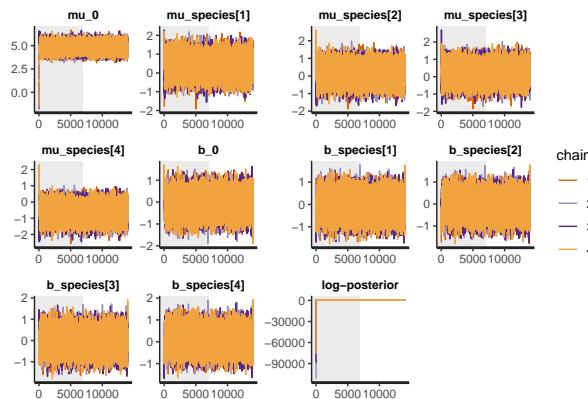
Before interpreting the posterior distributions of the parameters in this model, we assess the MCMC chains for convergence, the posterior distributions for normality, and the residuals of the model for normality and constant variance.

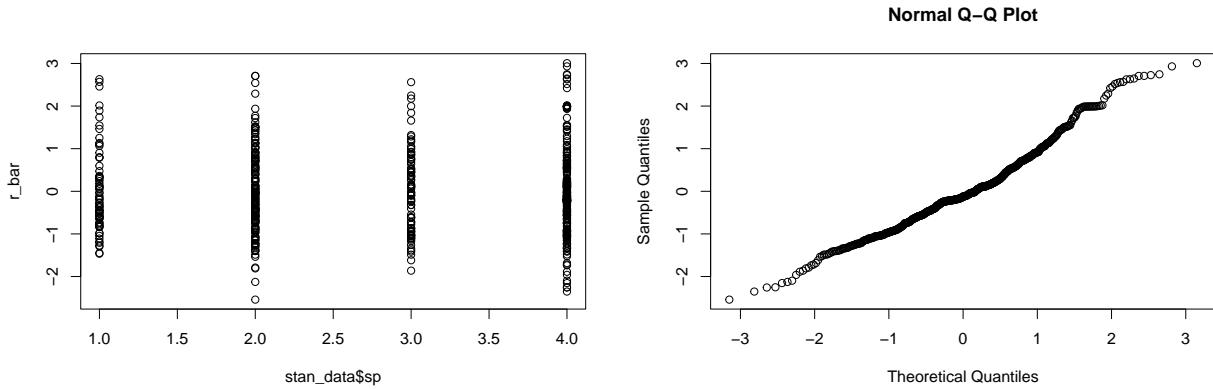
```
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=14000; warmup=7000; thin=1;
## post-warmup draws per chain=7000, total post-warmup draws=28000.
##
##           mean   se_mean     sd    10%    50%    90% n_eff Rhat
## mu_0      4.83    0.00  0.45   4.24   4.83   5.40  9915 1.00
## mu_species[1] 0.37    0.00  0.47  -0.23   0.36   0.97 10555 1.00
## mu_species[2] 0.13    0.00  0.46  -0.46   0.13   0.72 10441 1.00
## mu_species[3] 0.19    0.00  0.47  -0.41   0.19   0.79 10194 1.00
## mu_species[4] -0.66    0.00  0.45  -1.23  -0.66  -0.07  9744 1.00
## b_0      -0.08    0.00  0.44  -0.65  -0.08   0.49  9411 1.00
## b_species[1] -0.07    0.00  0.45  -0.64  -0.07   0.51  9579 1.00
## b_species[2] -0.08    0.00  0.45  -0.65  -0.07   0.49  9411 1.00
```

```

## b_species[3] -0.03 0.00 0.45 -0.61 -0.03 0.54 9566 1.00
## b_species[4] 0.09 0.00 0.44 -0.48 0.10 0.66 9398 1.00
## sigma_e[1] 0.76 0.00 0.07 0.68 0.75 0.84 40460 1.00
## sigma_e[2] 0.60 0.00 0.03 0.55 0.59 0.64 45397 1.00
## sigma_e[3] 0.79 0.00 0.06 0.71 0.79 0.87 40206 1.00
## sigma_e[4] 0.21 0.00 0.01 0.19 0.20 0.22 47928 1.00
## lp__ 792.07 2.25 28.08 764.88 785.08 829.02 156 1.03
##
## Samples were drawn using NUTS(diag_e) at Wed Aug 18 17:44:35 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```





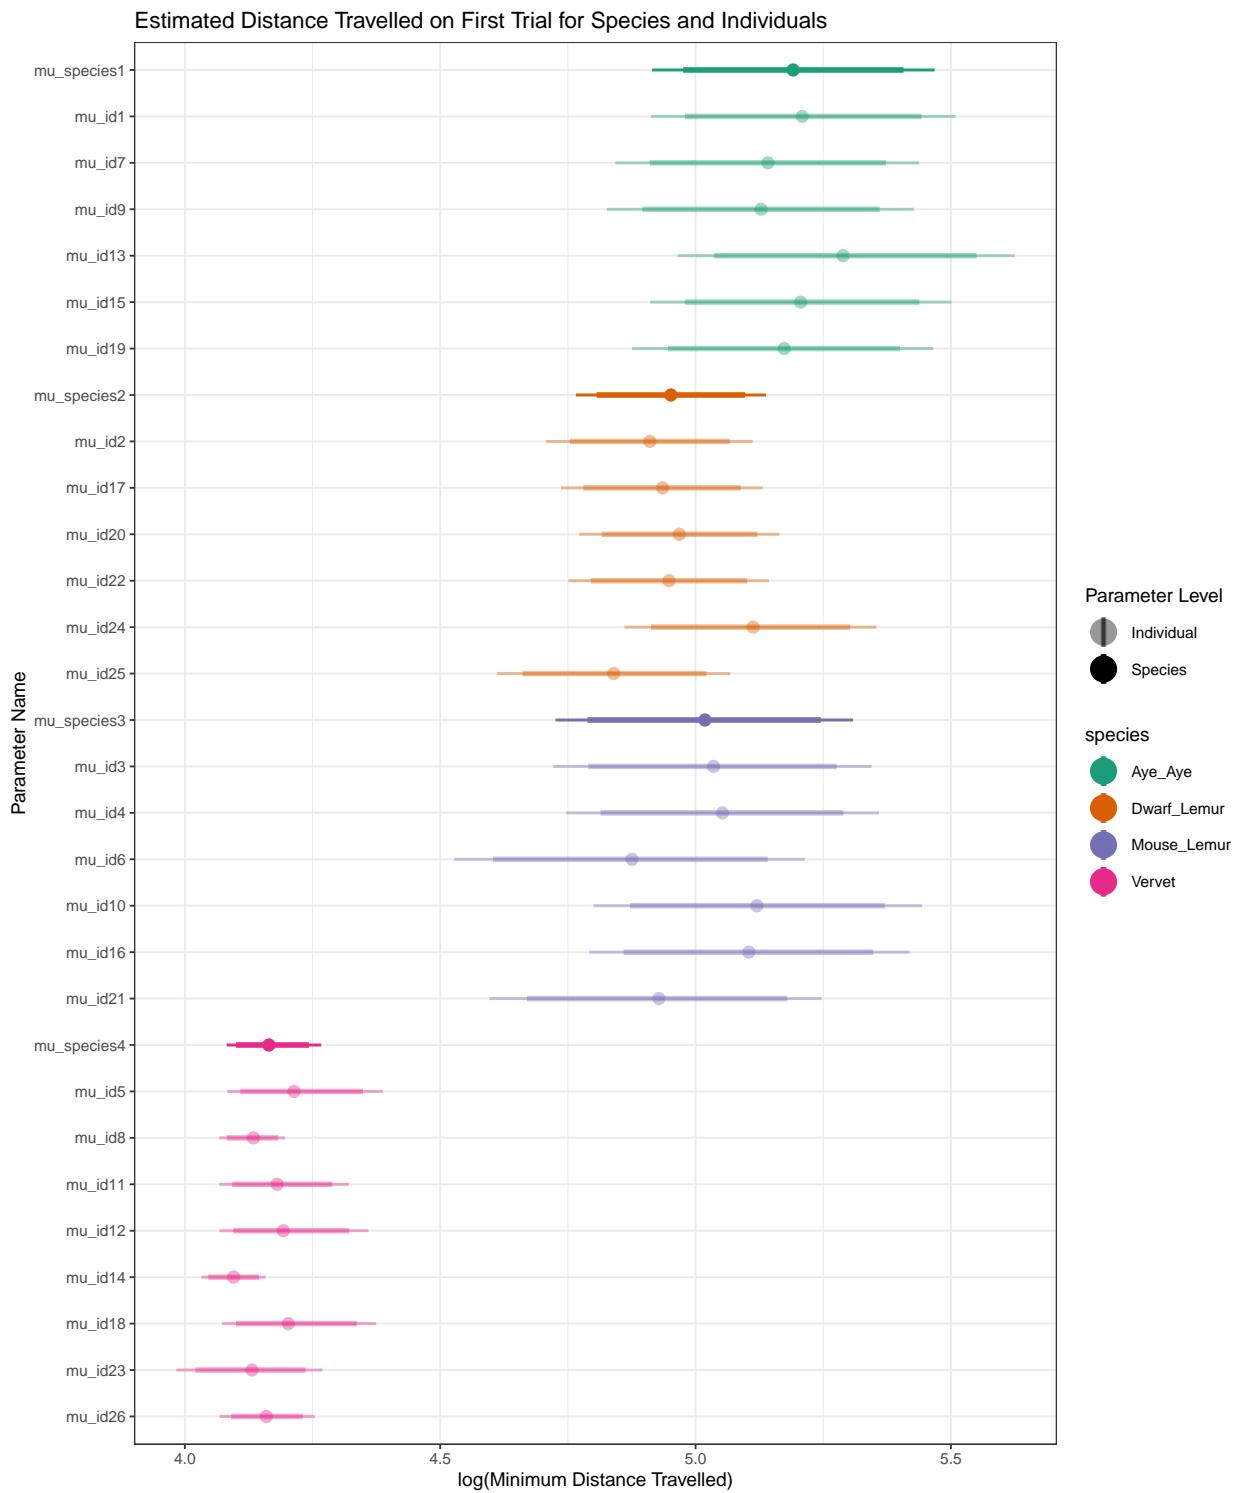
From the R-hats and effective sample sizes (`n_effs`) in the model summary (and from the visualization of parameter estimates across each MCMC chain) we can see that the posterior distribution of models converged across the four chains, and that the parameter space was well mixed in all chains. In pair plots of the model parameters, we can see that posterior distributions of all parameters appear normal, and there are no patterns to suggest the model was not identifiable. We do see high correlation among parameter estimates of the intercept and among parameter estimates of the slope. This is to be expected as we estimate both a grand mean for both parameters, and species level adjustments (normally distributed around zero); this means that as the estimated grand mean moves around in the posterior, species level adjustments will change accordingly to bring the species level intercept estimate (grand mean plus species adjustment) close to the true value.

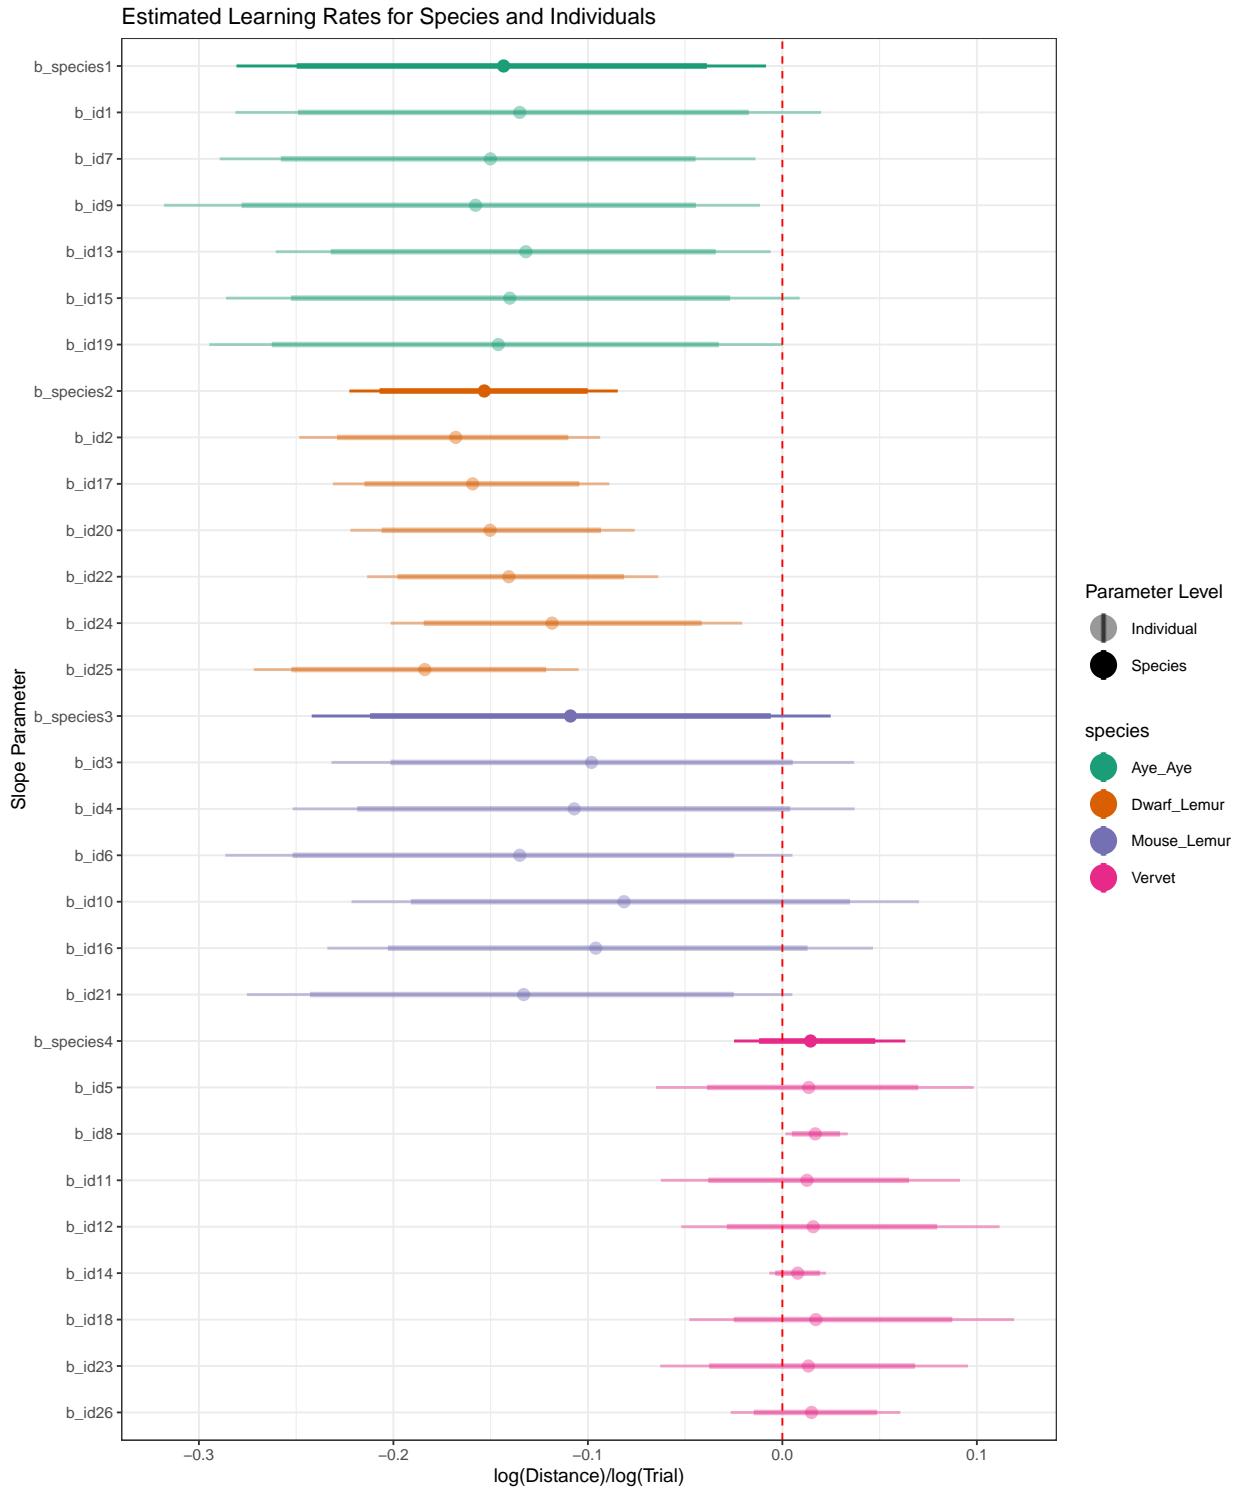
Plots of the residuals attest to the good fit between our statistical model and the structure of the data. Across species, trial number, and outcomes (distance traveled), residuals show no visible trend and remain evenly distributed.

Model Interpretation

Before assessing our posterior distributions, we need to add the estimated species and individual deviations from each sample of the posterior to the grand means from the same sample to find the actual posterior estimates of the intercept and slope for each species/individual.

To visualize the resulting posteriors, we determine quantiles of each posterior distribution and plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.



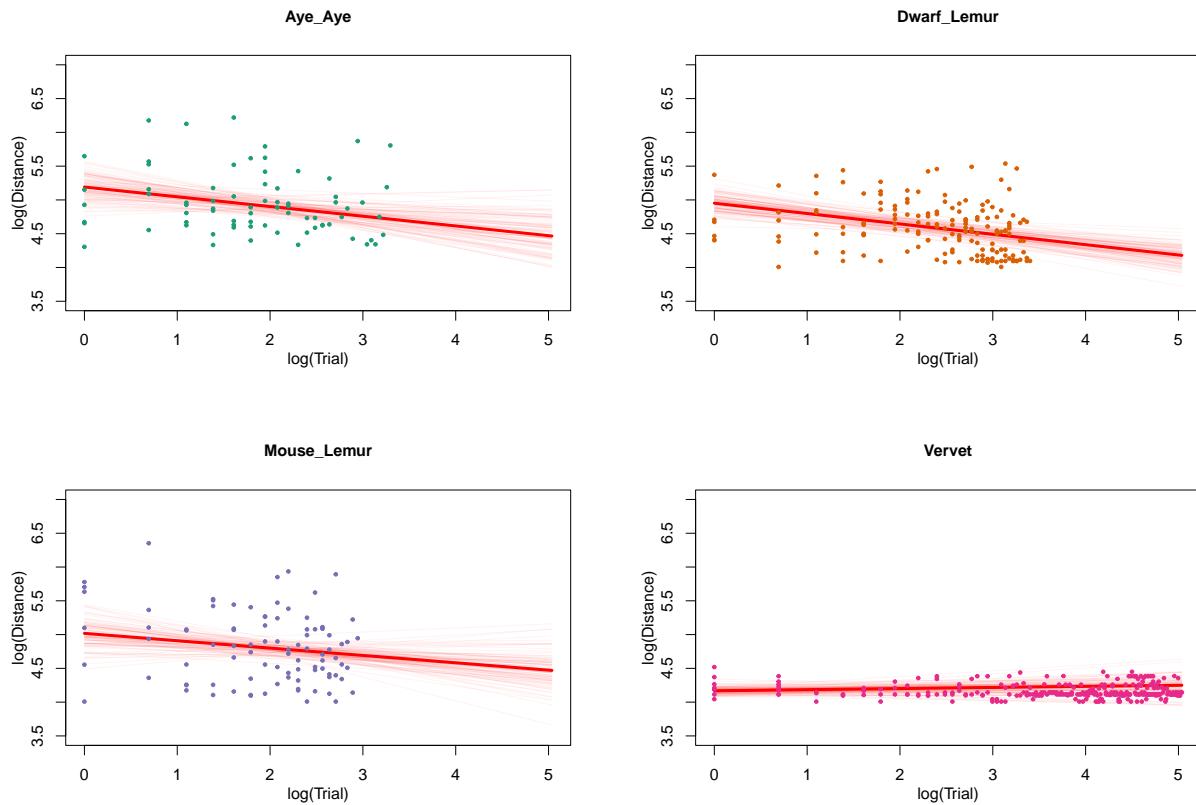


```
## pdf
## 2
```

```
## pdf
## 2
```

To visualize how the full set of parameter estimates come to together in our regression model, we plot a line using the mean intercept and slope for each species through the raw data. Additionally, we add the regression lines from 100 randomly selected samples using a more transparent color to illustrate the uncertainty of our estimates.

```
## pdf
## 2
```

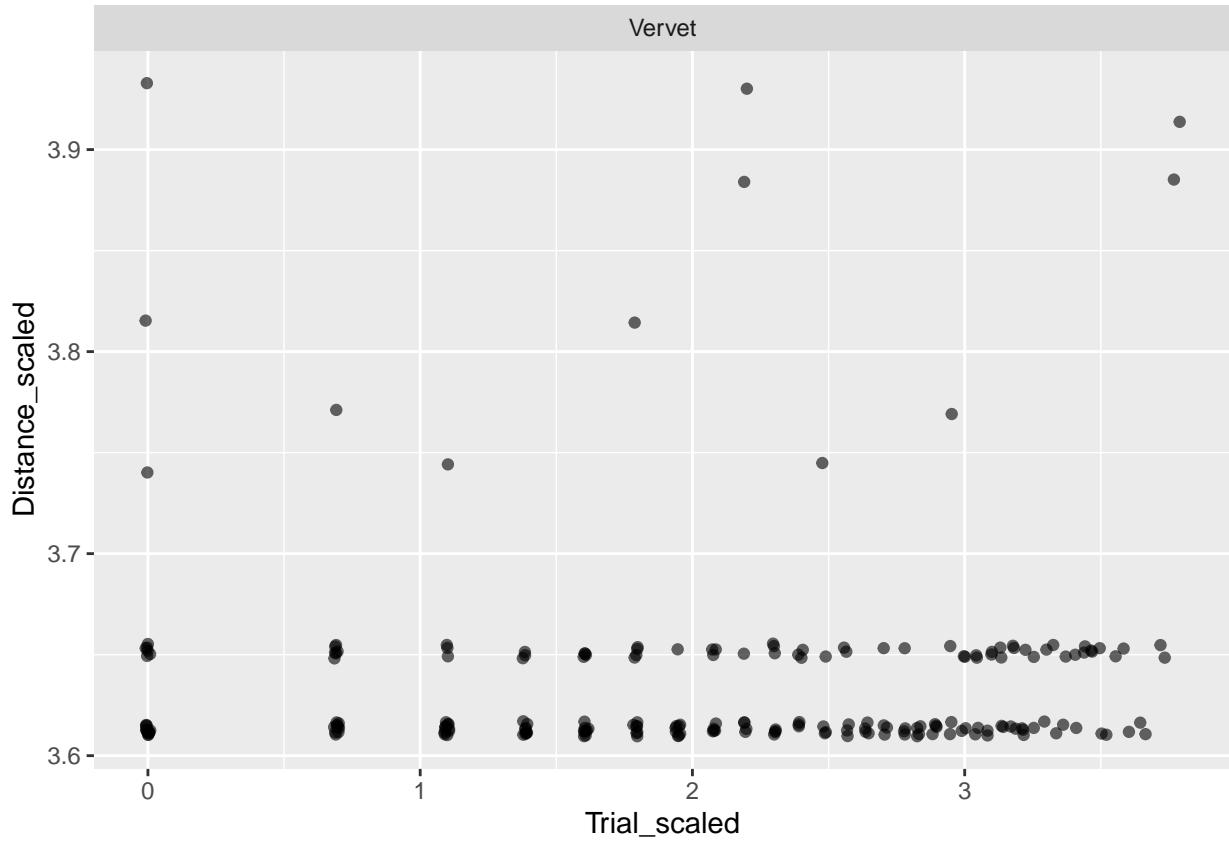


Finally, we report some specific numbers for use in final publication

```
## [1] 0.2043214
```

Pentagon Array Analysis

We analyze data from the Pentagon Array, which was only completed by vervets, in the same manner as the Double Trapezoid

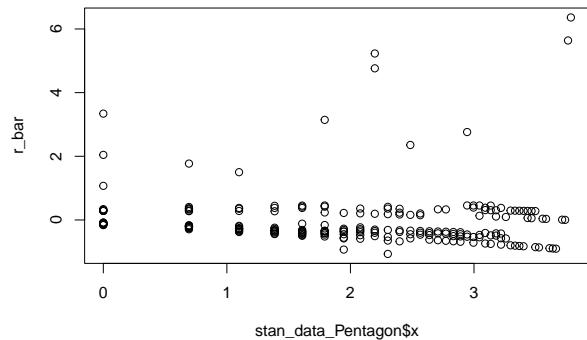
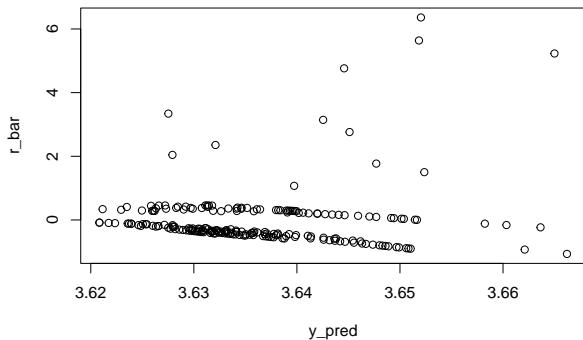
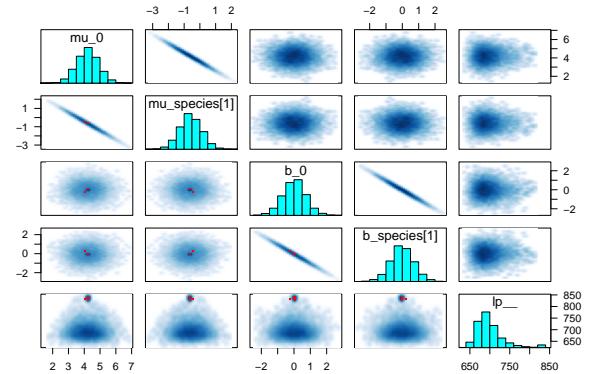
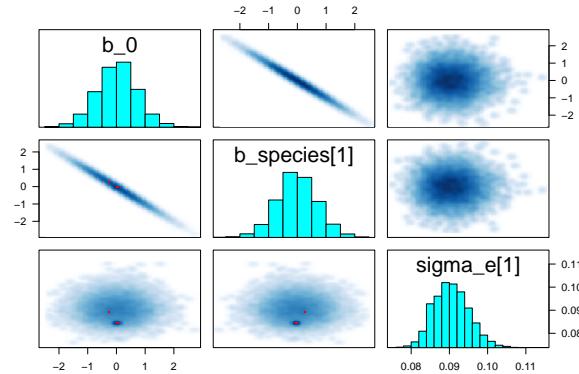
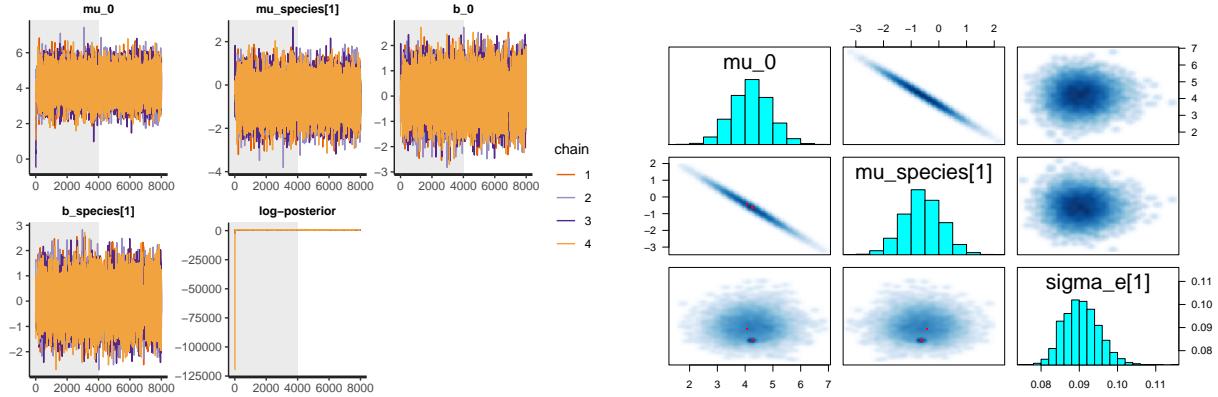


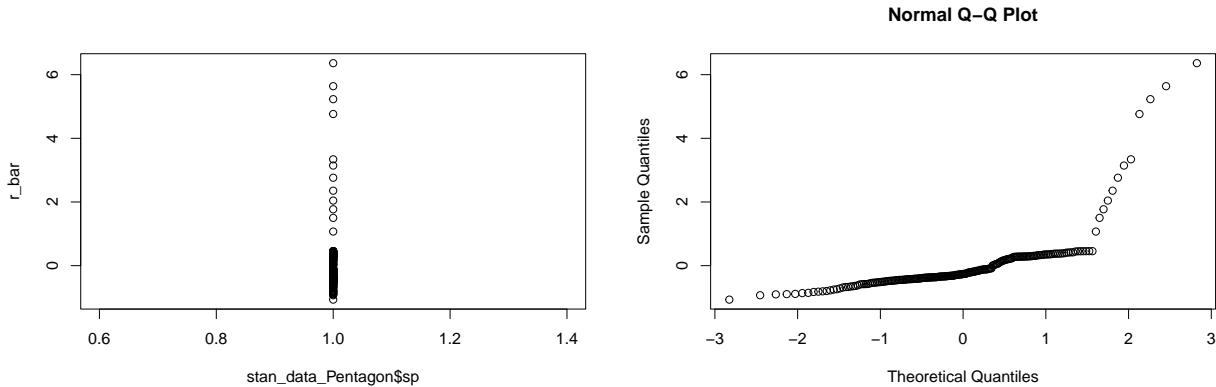
It can be seen from the plot of the raw data that vervets initially chose routes through the Pentagon array that were close to the shortest possible route, and that performance did not change much over time and experience. To analyze these patterns, we fit the data to the same statistical model as in the Double Trapezoid.

Model Diagnostics

After fitting the model in STAN, we assess the resulting MCMC chains and posterior distributions for convergence and mixing, then examine the residuals to assess how well the model fits the structure of the data

```
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=8000; warmup=4000; thin=1;
## post-warmup draws per chain=4000, total post-warmup draws=16000.
##
##           mean   se_mean     sd    10%    50%    90% n_eff Rhat
## mu_0       4.21     0.01   0.69   3.32   4.24   5.10  8082 1.00
## mu_species[1] -0.59     0.01   0.69  -1.47  -0.61   0.30  8094 1.00
## b_0        0.01     0.01   0.70  -0.88   0.03   0.89  7008 1.00
## b_species[1]  0.00     0.01   0.70  -0.88  -0.02   0.89  7006 1.00
## sigma_e[1]    0.09     0.00   0.00    0.08   0.09   0.10 1223  1.00
## lp__      699.84    4.85 35.16 666.20 692.34 739.93    53 1.08
##
## Samples were drawn using NUTS(diag_e) at Tue Aug 24 14:29:13 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```





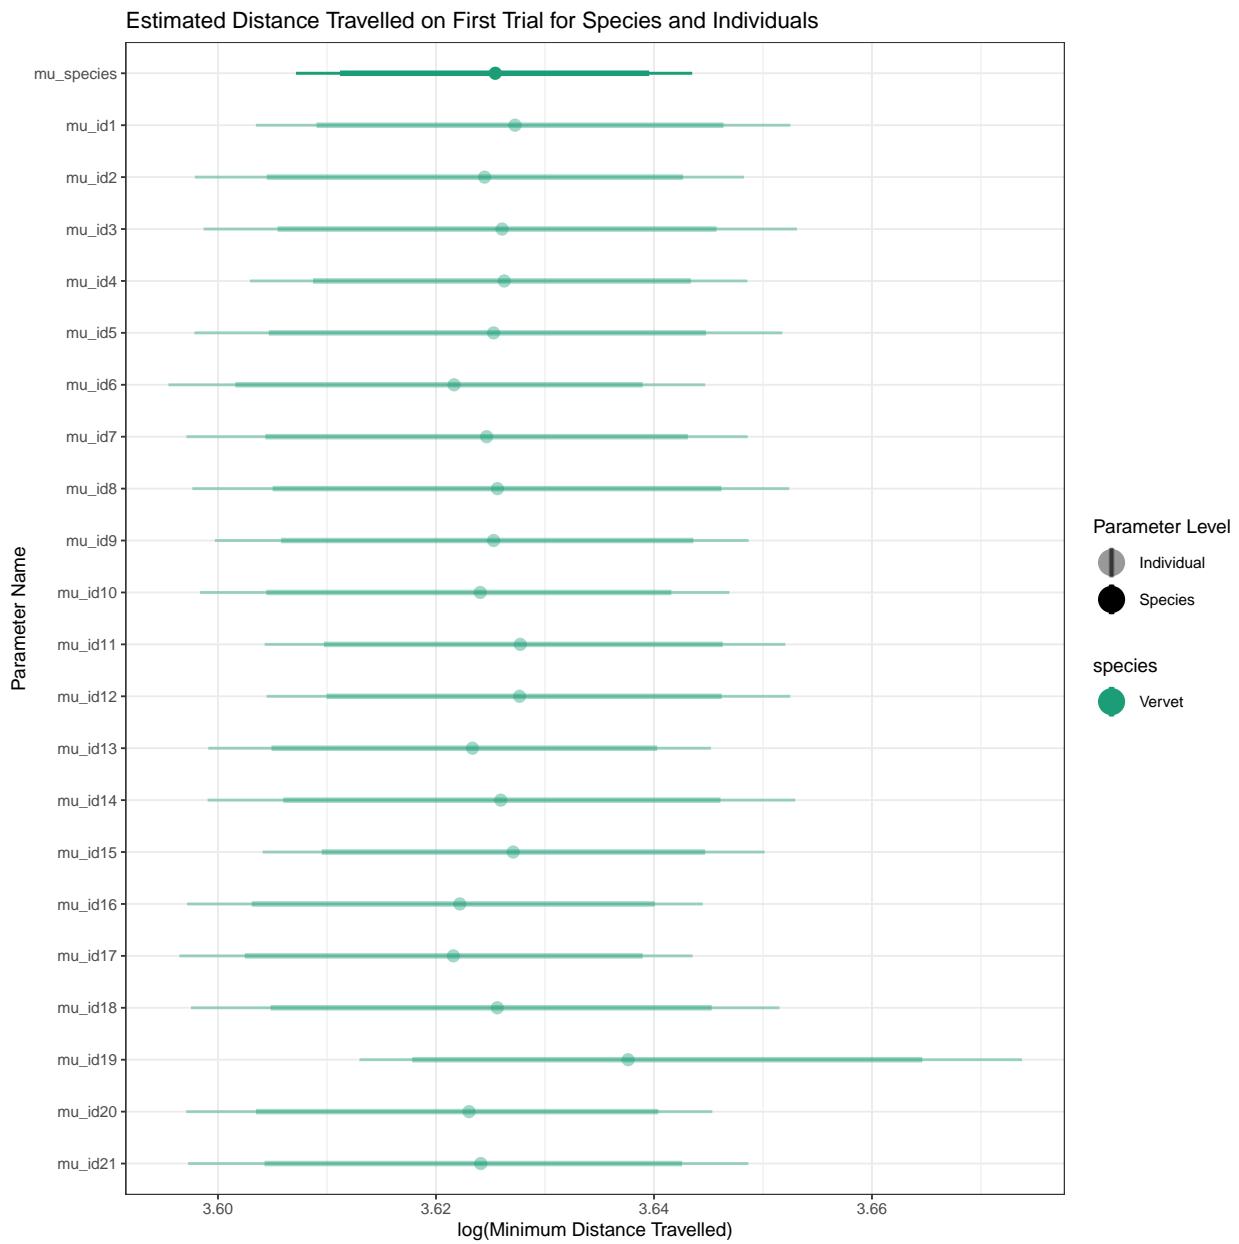
The MCM chains converge and show good mixing. The pairs plots indicate that parameter distributions in the posterior are normally distributed, with some expected correlation structure as described in the double trapezoid section. Red pixels in the pairs plot indicate samples with divergent transitions. Their placement in the σ_e distributions suggest that these have to do with challenges estimating the correct error variance, while their placement in the distribution of b_0 suggests that this only occurs when the slope is estimated at 0. Their placement in the distribution of log-probability of the samples (high) suggests that when the model slope is estimated at 0, the error variance can be reduced, producing a better fit model, but that this requires a large change in the estimated error variance that creates warnings. Taken together, these divergent transitions highlight that running a model with species level deviations around a grand mean for data with only 1 species and structure, discrete outcomes in a continuous variable isn't the best approach. We might consider simplifying the model for this array.

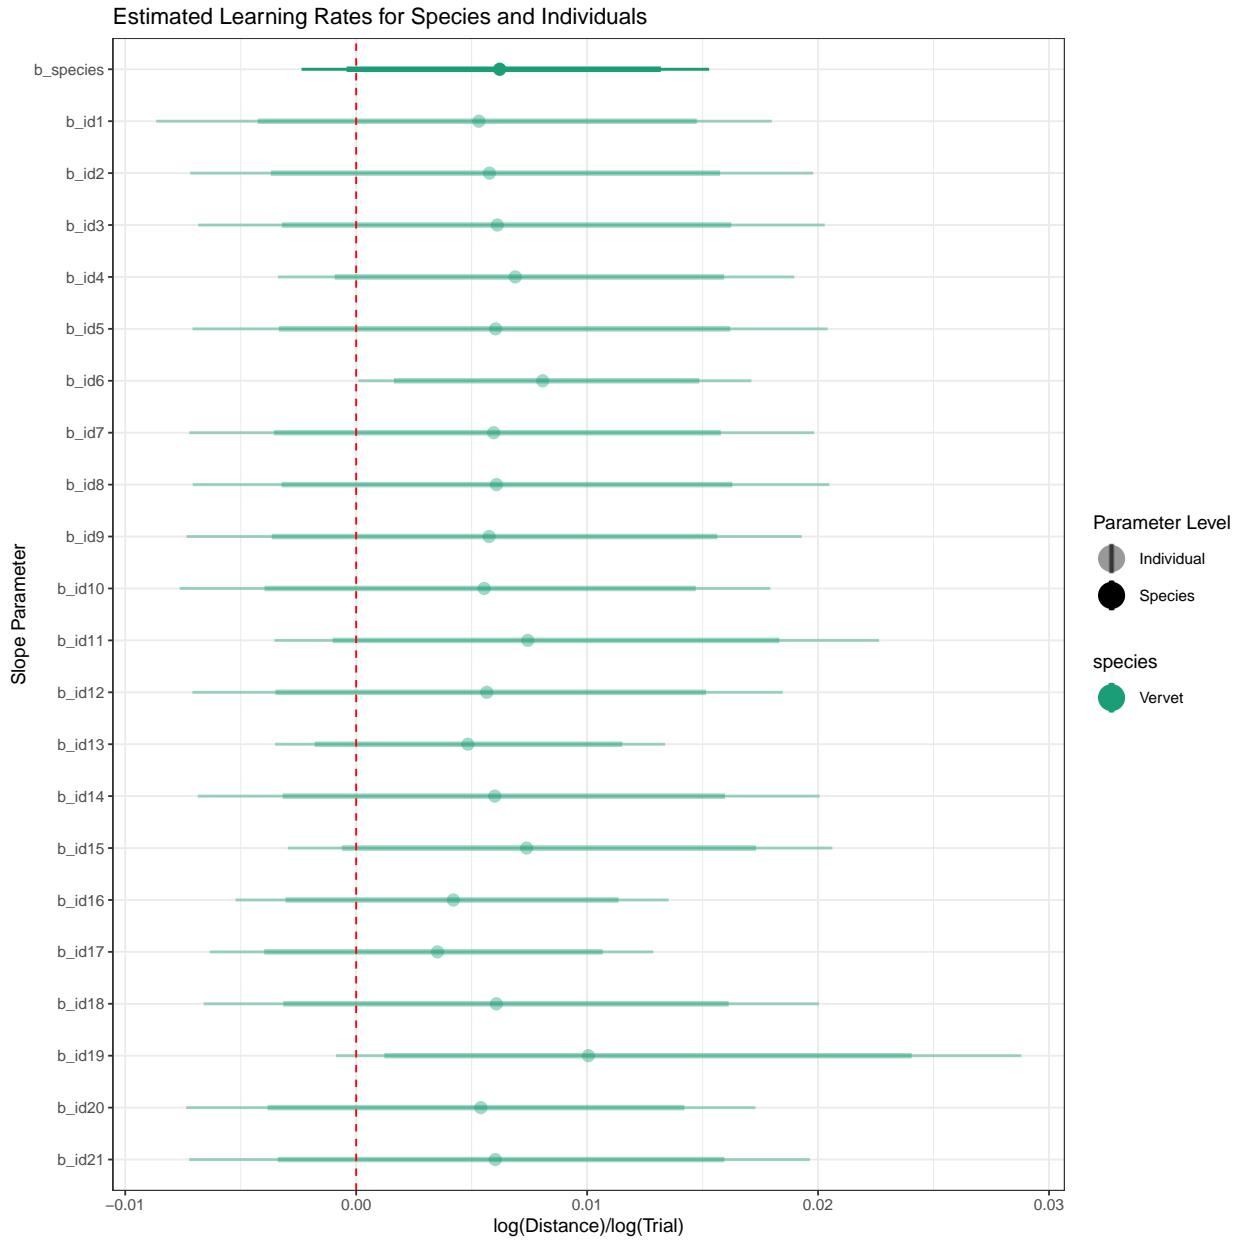
Accordingly, when looking at the residuals, it is revealed that this statistical model does not fit the structure of the data quite as well as in the Double Trapezoid. Based on the patterns in the residual plots and what we know about the structure of this experiment, this is likely because the outcome (distance traveled) has little room for gradual improvement: instead vervets are switching between a limited set of options with two, discrete possibilities for outcome, and then occasionally switching to other, less effective routes. Likely, these data would be better modeled by some sort of discrete choice following a Poisson distribution. However, to enable comparisons between arrays, we stick with this model for analysis, and make note of the inconsistencies as a caveat in interpreting our results.

Model Interpretation

Before assessing the posterior distributions, we need to add the deviations for species and individuals to the grand means for each sample to find the actual posterior estimates of the intercept and slope for each species/individual.

To visualize the resulting posteriors, we determine quantiles of each posterior distribution and plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.



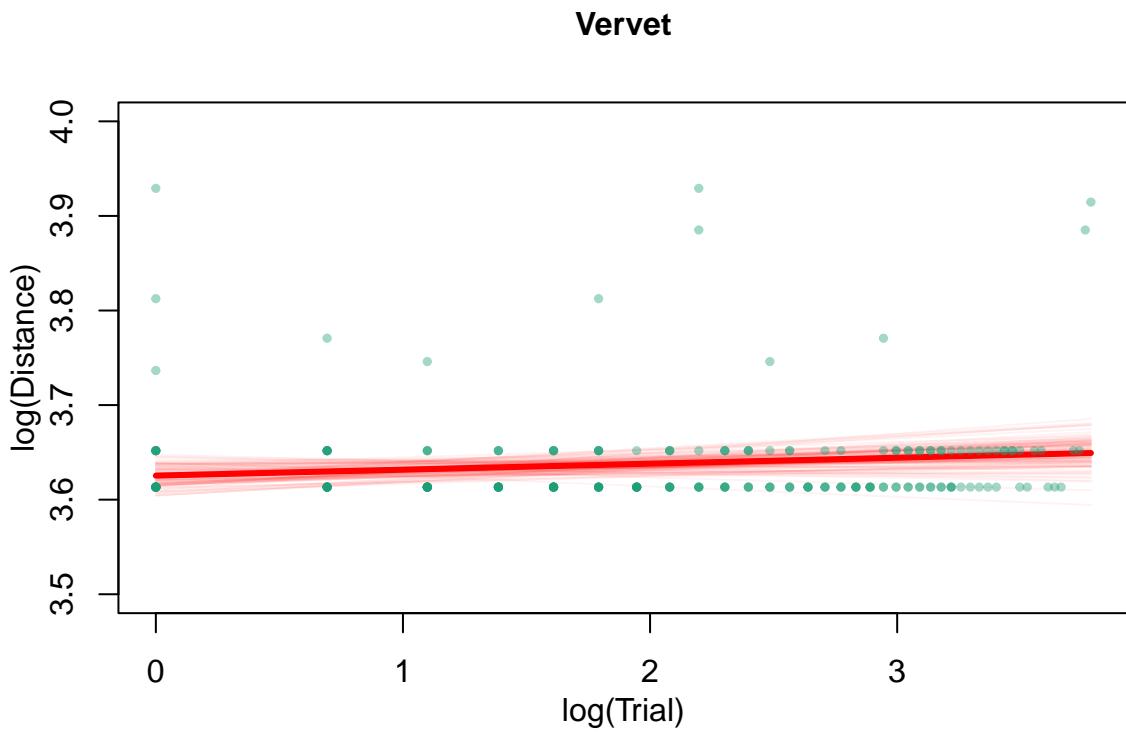


```
## pdf
## 2
```

```
## pdf
## 2
```

Finally, we plot a mean regression line (thick) and 100 sampled regression lines (thin) through the data for each species.

```
## pdf
## 2
```

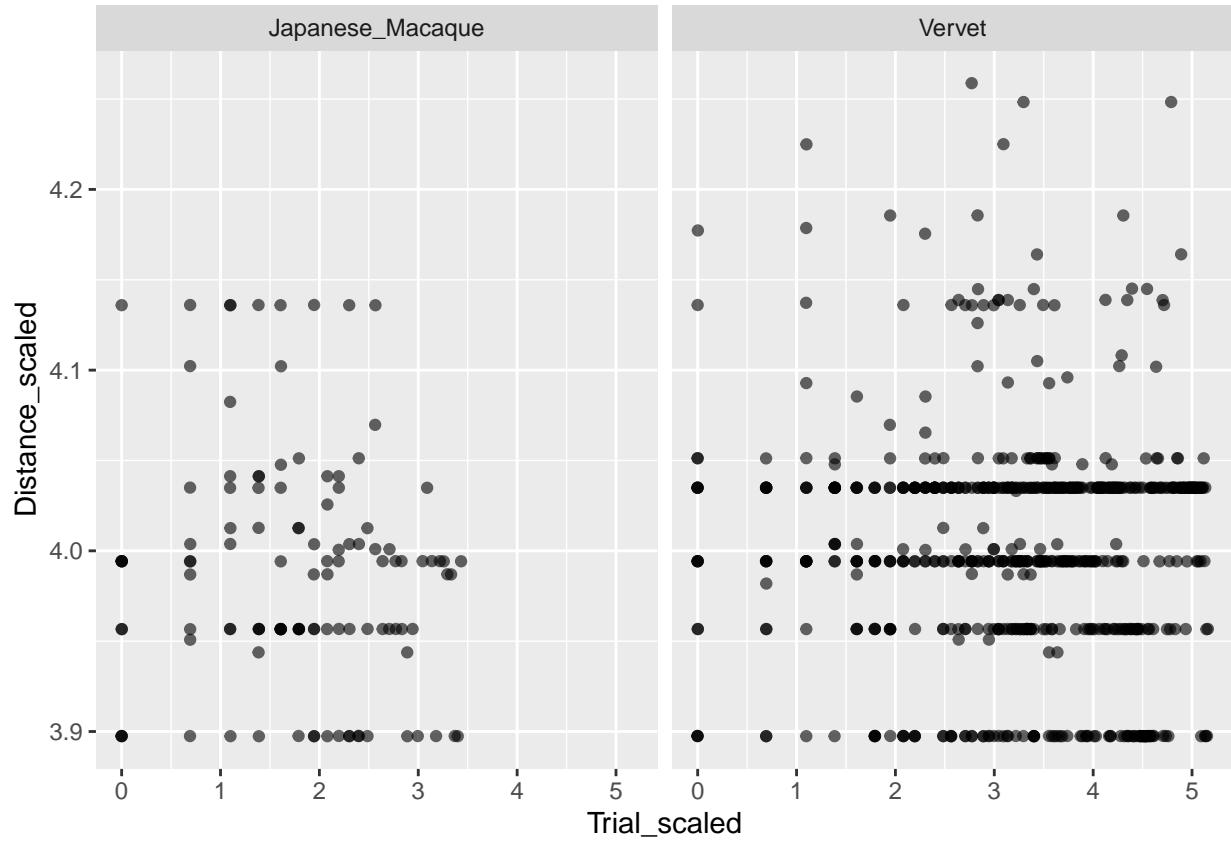


In the below code, we extract a few values for reporting in final publication

```
## [1] 0.885
```

Z-Array Analysis

We analyze data from the Z-Array, which was completed by vervets and japanese macaques, in the same manner as the Double Trapezoid and the Pentagon.

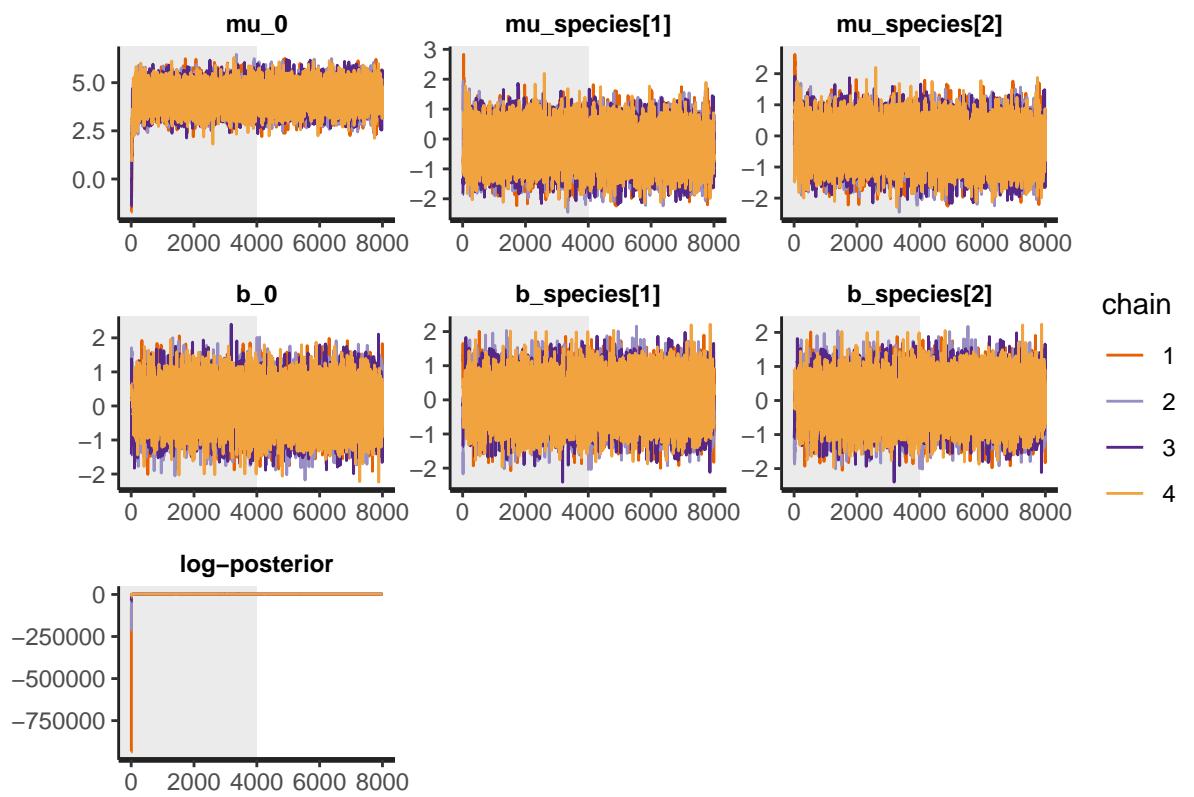


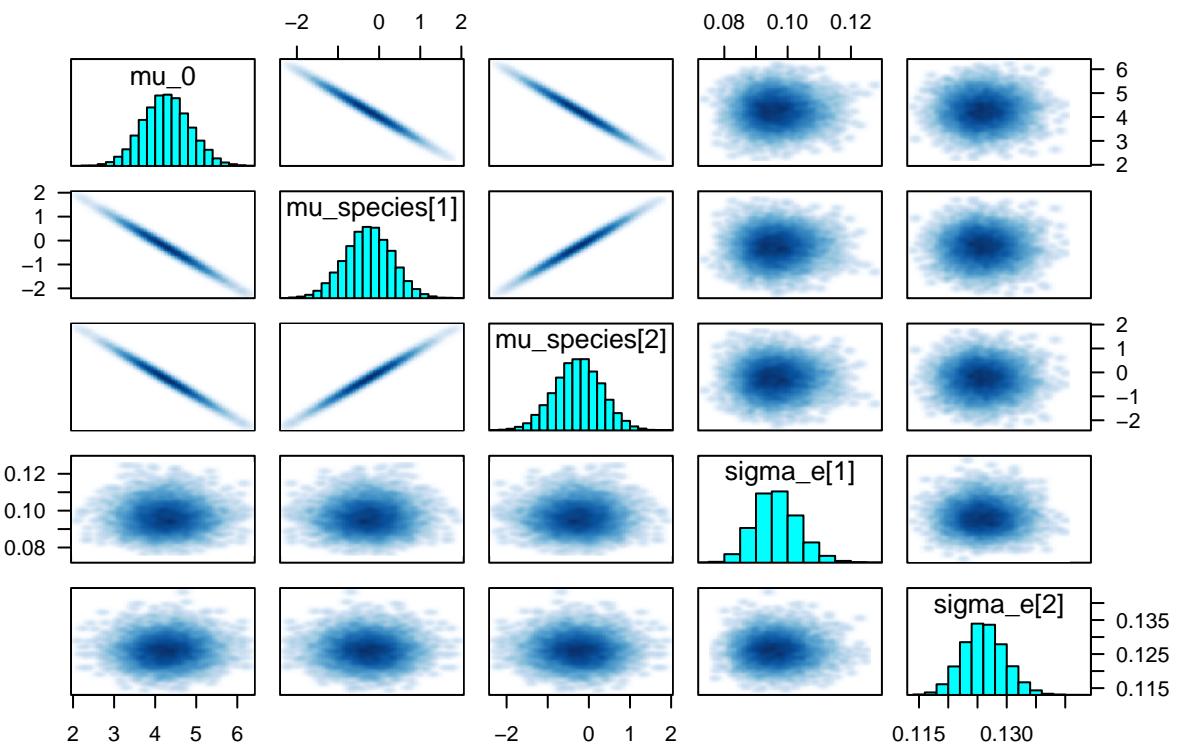
Model Diagnostics

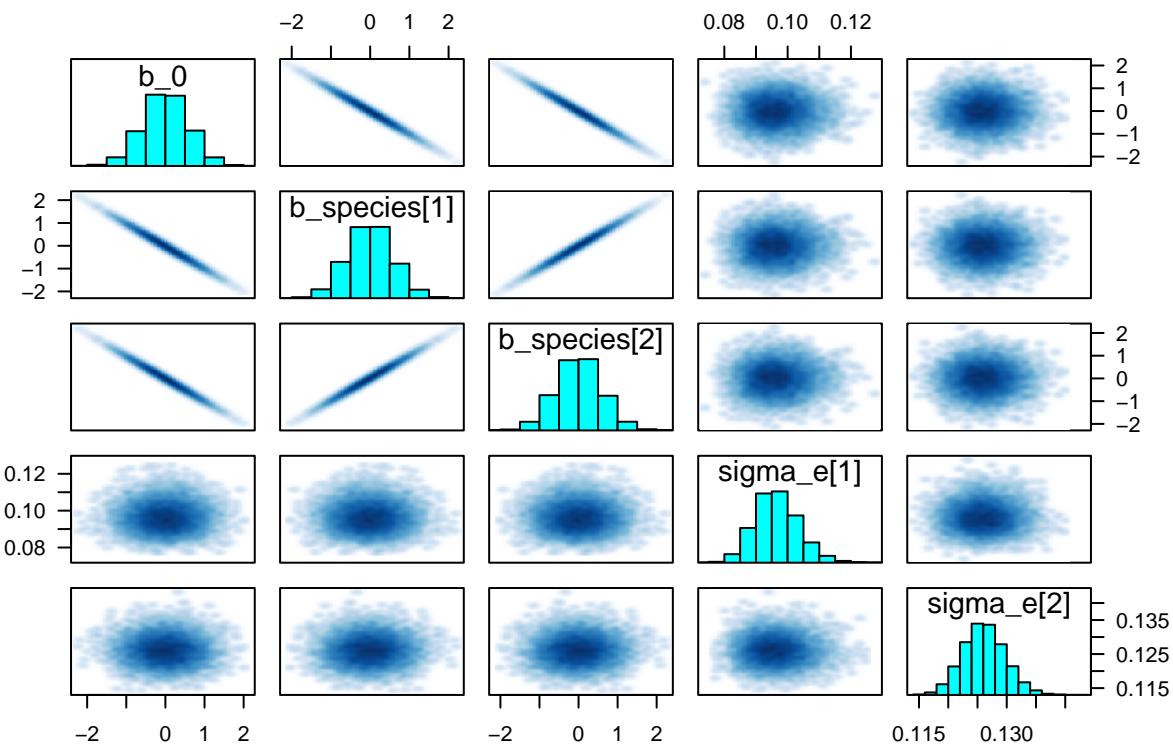
```

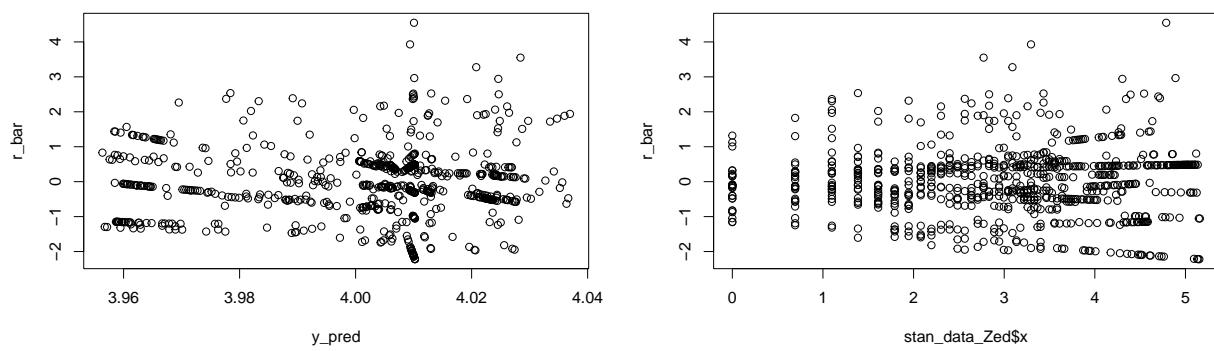
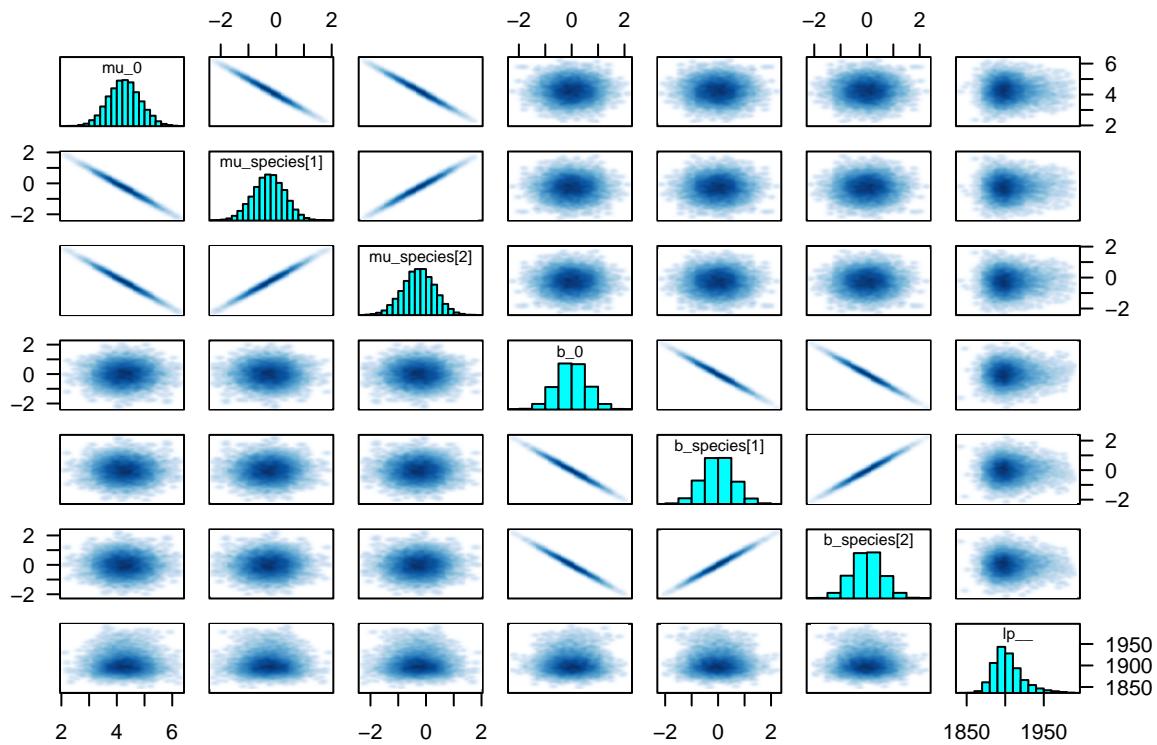
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=8000; warmup=4000; thin=1;
## post-warmup draws per chain=4000, total post-warmup draws=16000.
##
##          mean   se_mean     sd    10%    50%    90% n_eff Rhat
## mu_0      4.27    0.01  0.58   3.54   4.26   5.02  7398   1
## mu_species[1] -0.26    0.01  0.58  -1.01  -0.26   0.47  7396   1
## mu_species[2] -0.26    0.01  0.58  -1.01  -0.26   0.47  7391   1
## b_0       0.00    0.01  0.58  -0.74   0.00   0.74  6861   1
## b_species[1] -0.01    0.01  0.58  -0.75  -0.01   0.73  6862   1
## b_species[2]  0.00    0.01  0.58  -0.74   0.00   0.74  6857   1
## sigma_e[1]   0.10    0.00  0.01   0.09   0.10   0.11 17021   1
## sigma_e[2]   0.13    0.00  0.00   0.12   0.13   0.13 20737   1
## lp__      1903.85  1.06 19.29 1882.86 1900.72 1929.61    334   1
##
## Samples were drawn using NUTS(diag_e) at Thu Aug 26 00:50:36 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

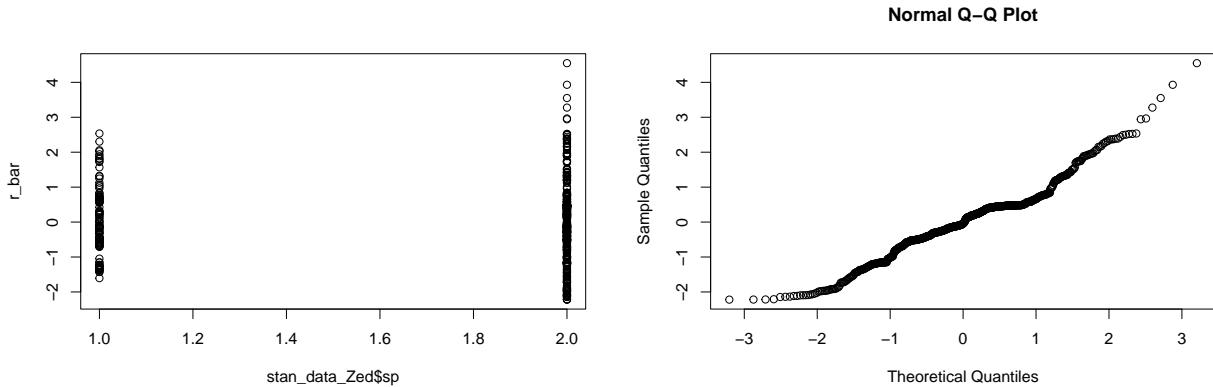
```









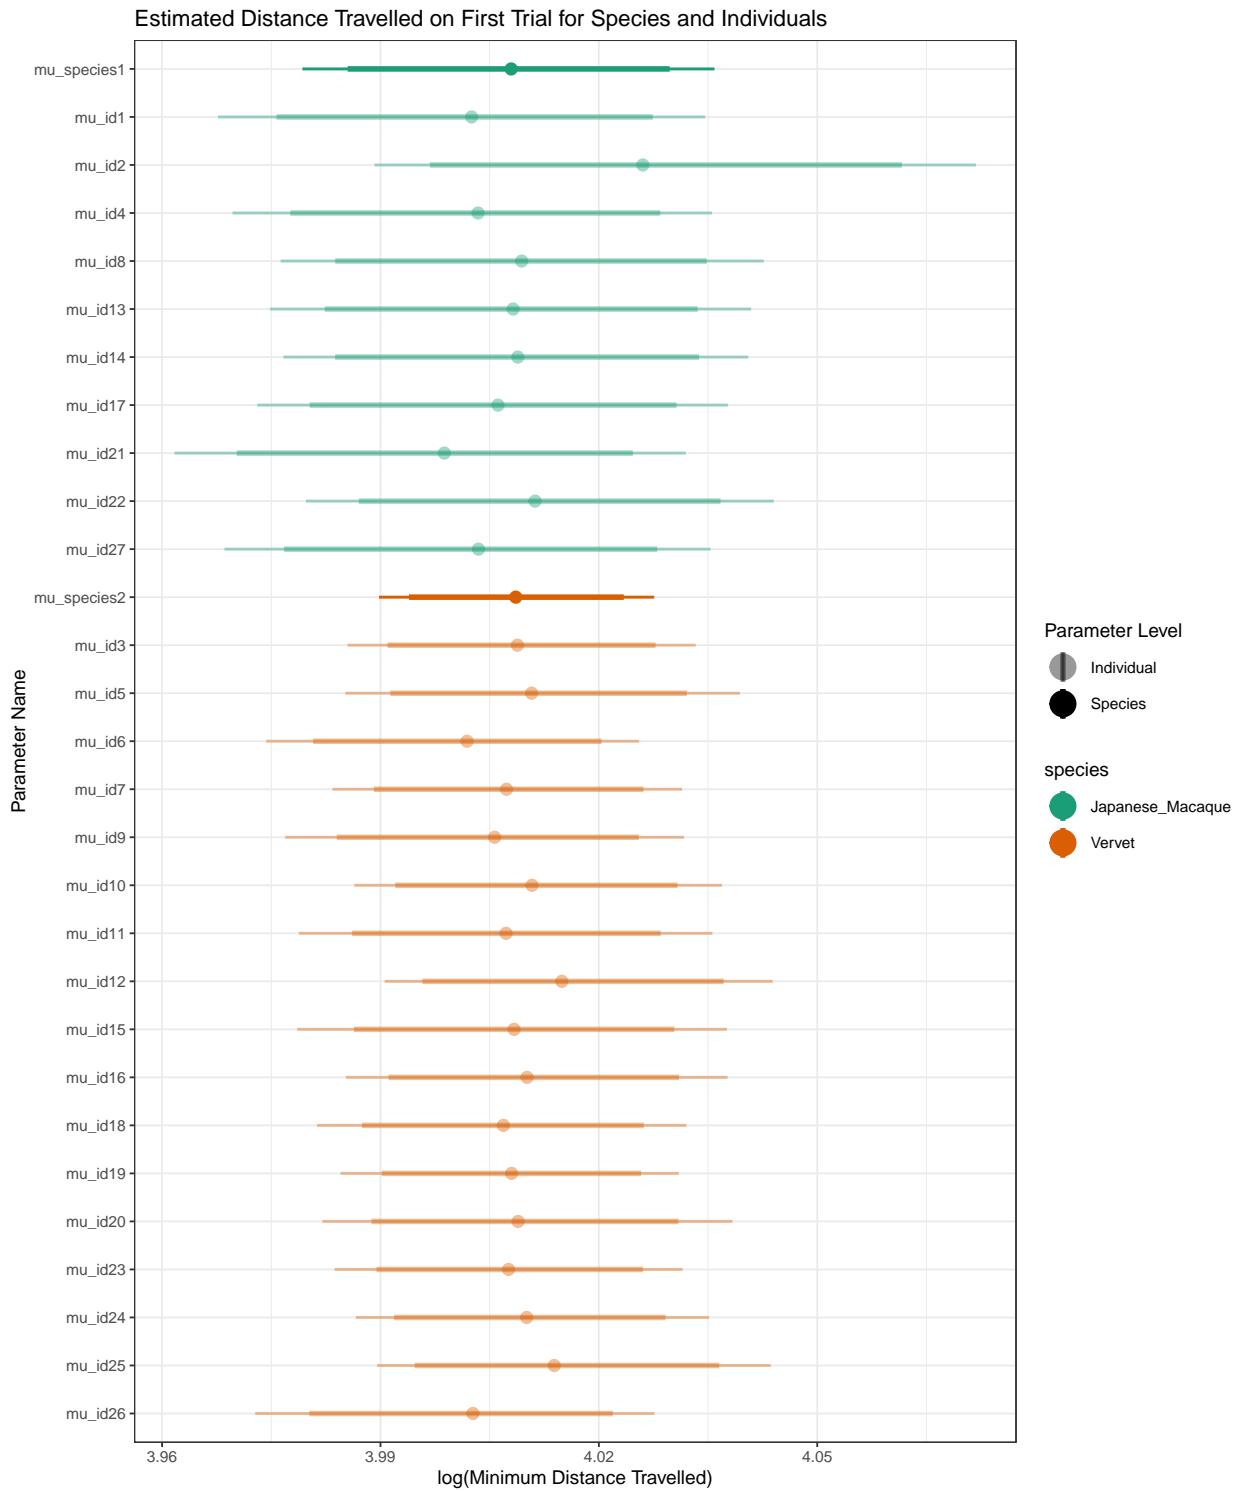


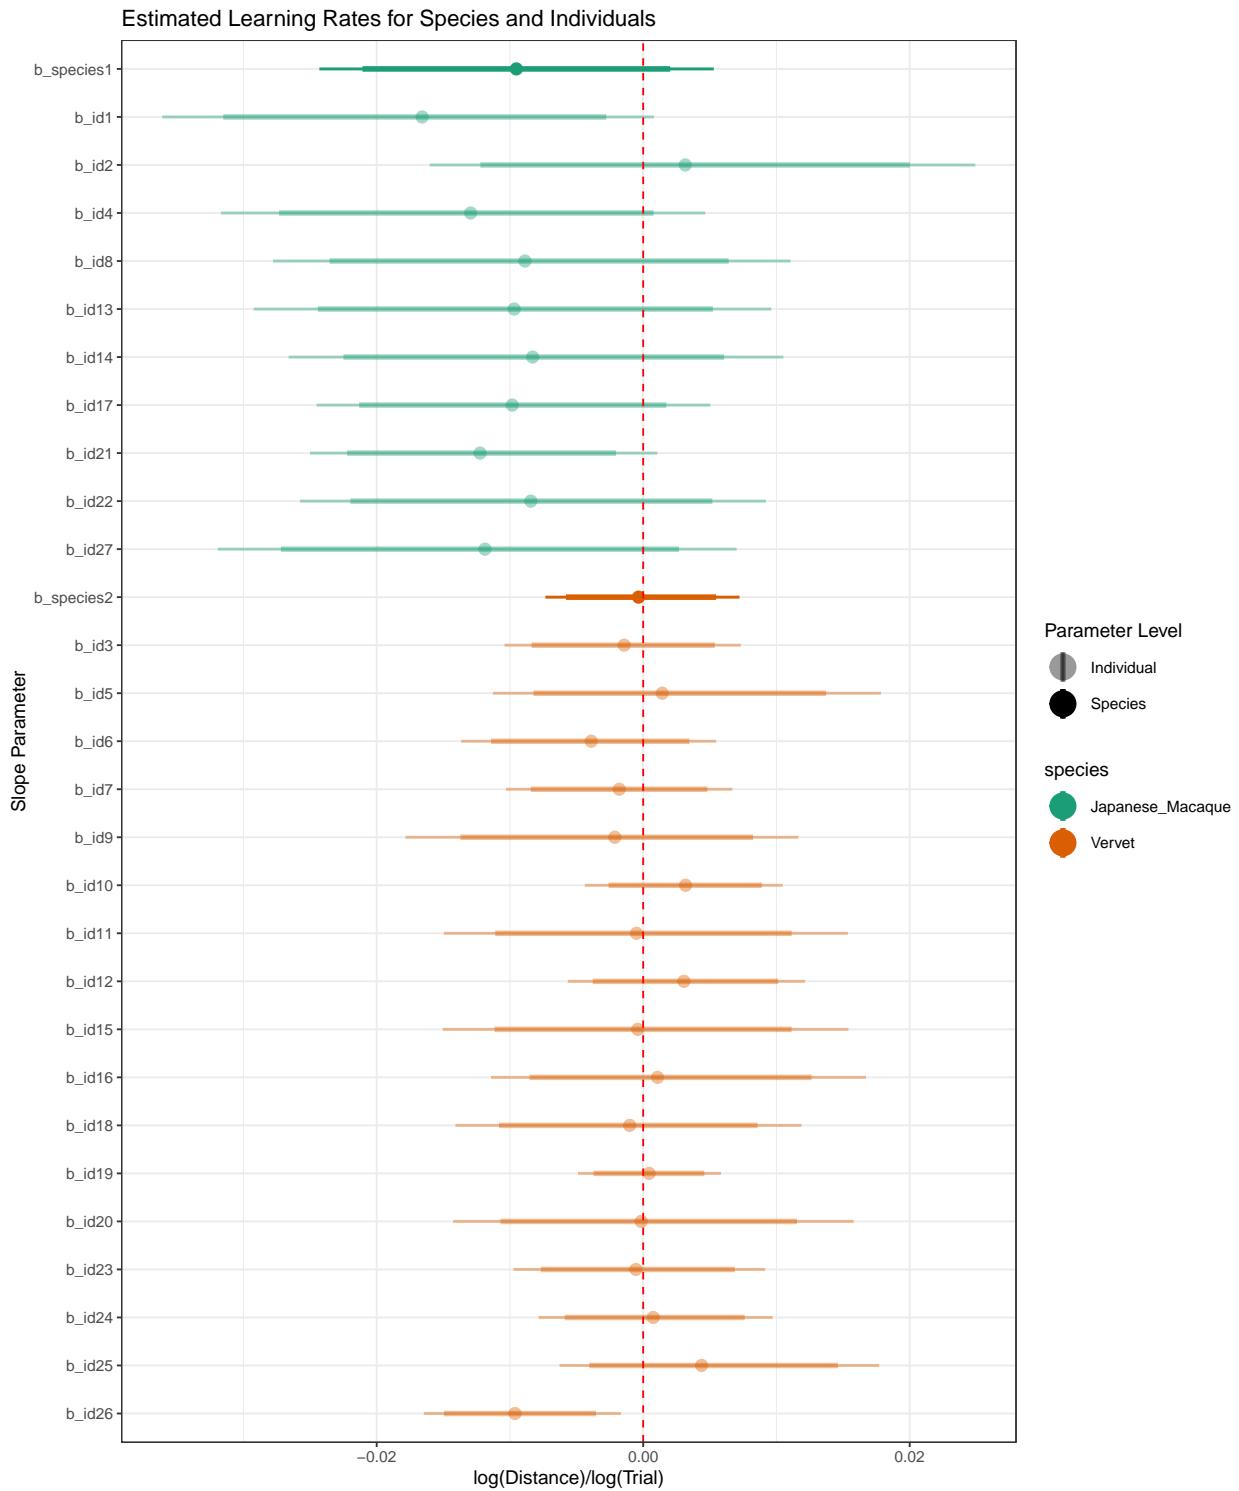
MCMC chains show converging and good mixing. Parameter distributions in the posterior are normal with no unexpected correlations. Residuals show a slight downward trend with predicted distance traveled, but it is small enough not to be concerning.

Model Interpretation

Because we are estimating the grand mean for the intercept and slope of the model and then estimating species and individual level *deviations*, we need to add the deviations to the grand means for each sample to find the actual posterior estimates of the intercept and slope for each species/individual.

To visualize the resulting posteriors, we determine quantiles of each posterior distribution and plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.



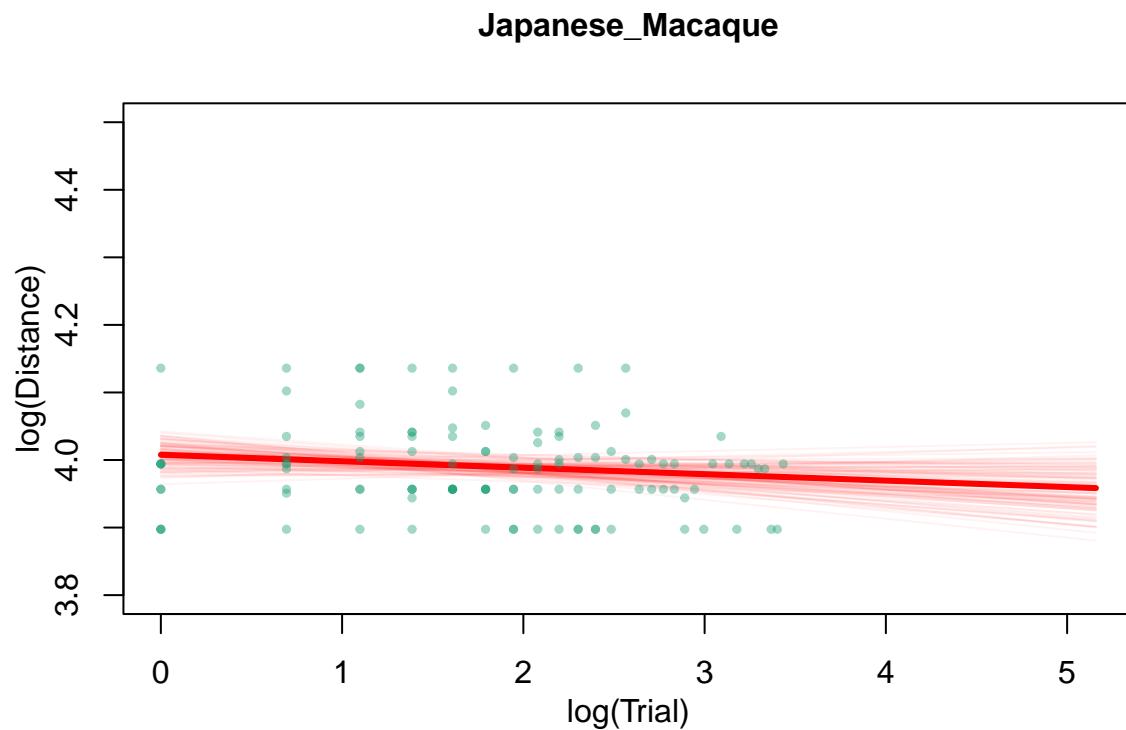


```
## pdf
## 2
```

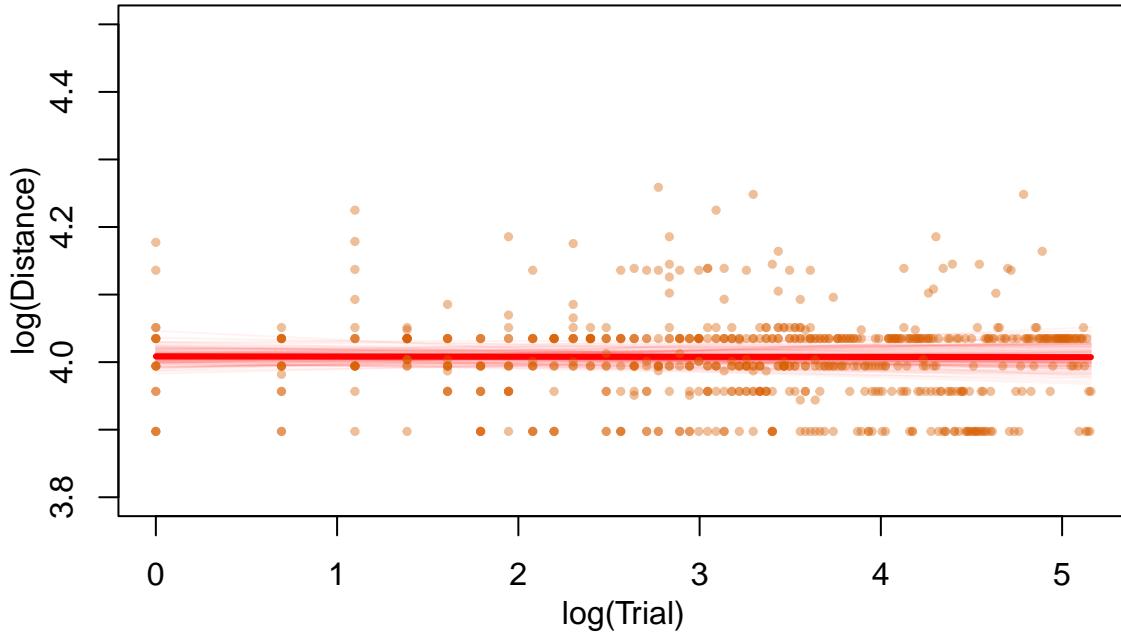
```
## pdf
## 2
```

Finally, we plot a mean regression line (thick) and 100 sampled regression lines (thin) through the data for each species.

```
## pdf  
## 2
```



Vervet



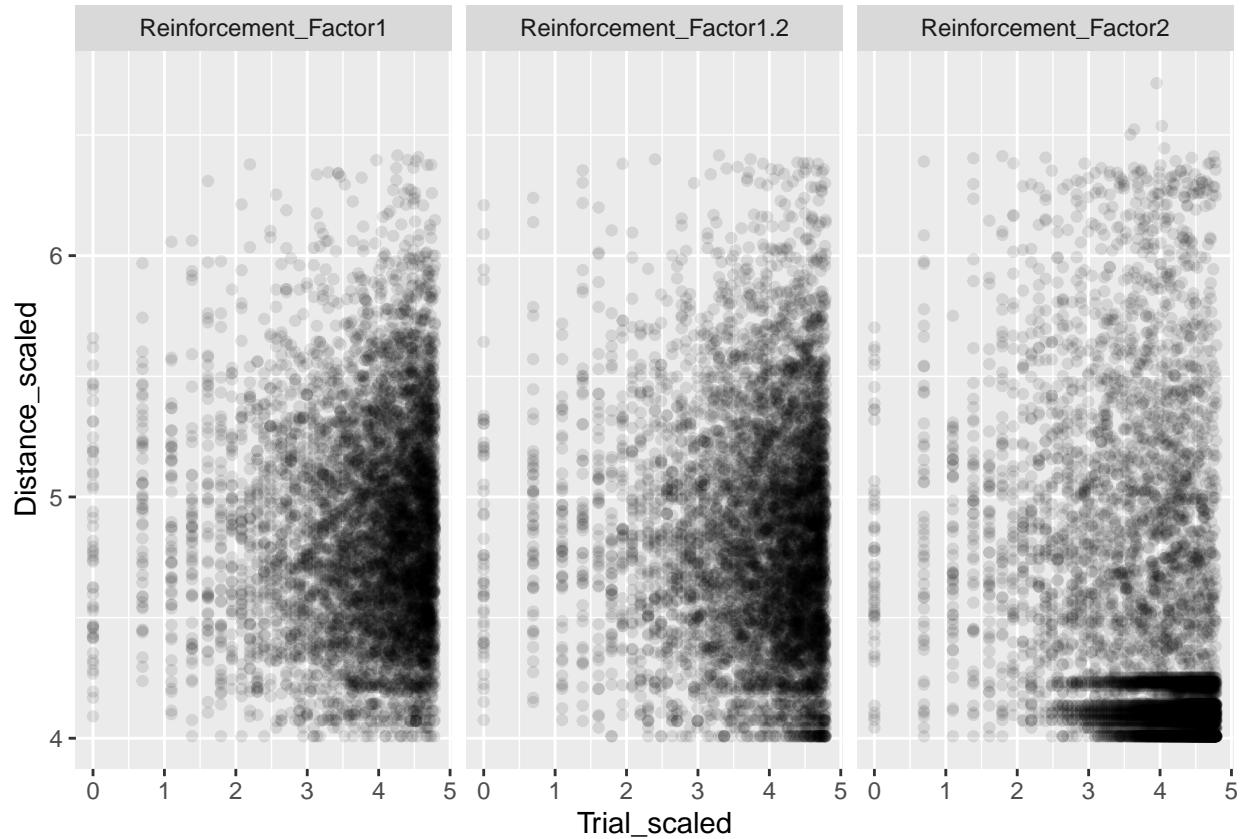
Below, we calculate a few additional values for final publication

```
## [1] 0.855
```

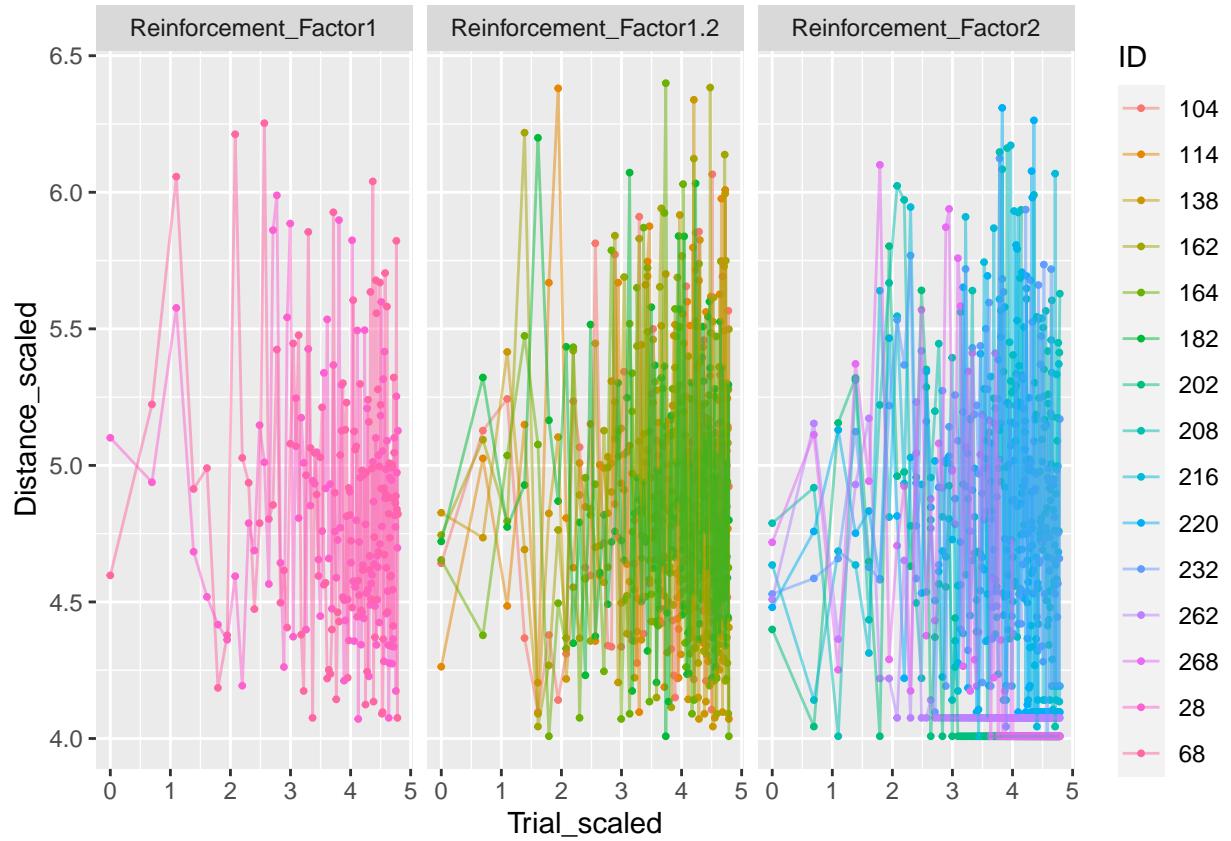
Simulation analysis, Double Trapezoid

We start by analyzing simulations from the Double Trapezoid array, in which the most species of primates were tested and thus offers the most potential for comparative analysis.

In our simulation model, an agent first moved between array locations using a predetermined transition matrix. Agents, like our primate subjects, completed a trial once all locations had been visited. Upon completing a trial, the agent compared the total distance traveled in that trial to a stored (remembered) shortest-distance-to-trial-completion. If the current trial was shorter than that known previously, the agent changed its transition matrix by multiplying the probability of all transitions that occurred in the current trial by a reinforcement factor (e.g., 1.2), then renormalized the matrix so that all transition probabilities from a given location sum to 1. The agent repeated this process until 120 trials were completed.



Before running the analysis, we can already see from the raw data that simulations with a Reinforcement Factor of 2 begin to converge of usage of routes with relatively low total travel distance around $e^{2.5}$ (~12) trials, though a subset of individuals remain highly varied in their performance. No such trend is immediately visible in data from simulations with Reinforcement Factors of 1 or 1.2. To better see the heterogeneity in individual performance, especially in Reinforcement Factor 2 agents, we add lines to the data for a subset of these individuals



Having gained a sense of the patterns in these simulations, we analyze them using the same statistical model as used for the empirical data. The final STAN properties used in the following code block we adjusted several times to achieve convergence and mixing in the MCMC chains.

Our first step after computing model parameters in STAN is to visualize diagnostics of well the model fit the data.

```

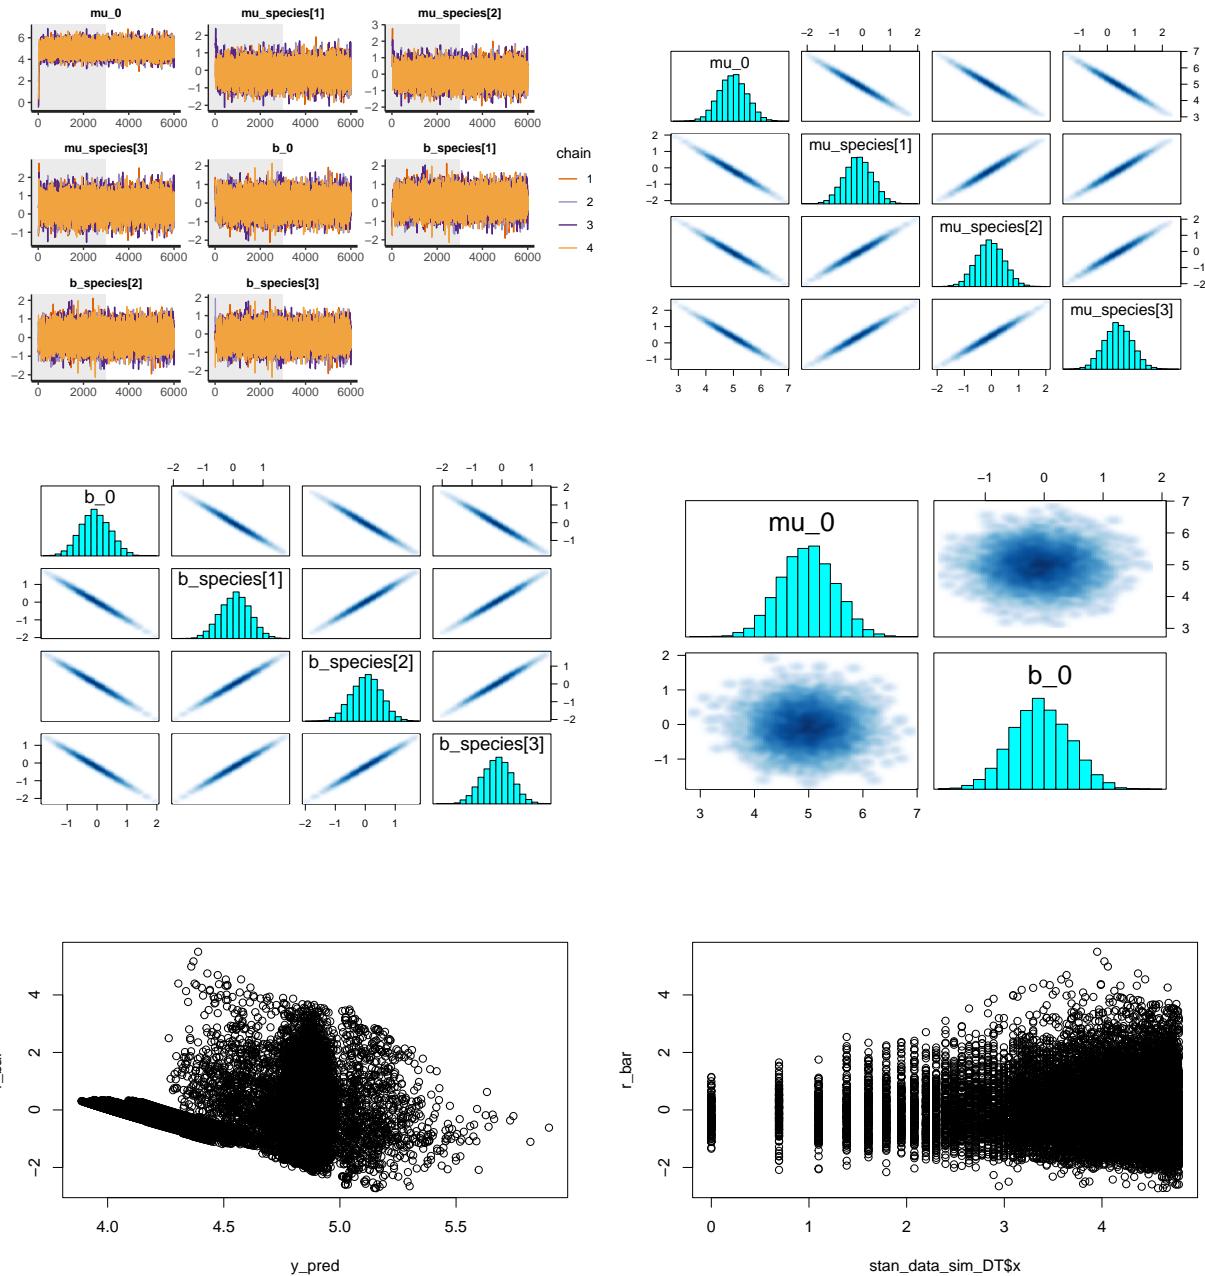
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=6000; warmup=3000; thin=1;
## post-warmup draws per chain=3000, total post-warmup draws=12000.
##
##           mean se_mean    sd    10%    50%    90% n_eff Rhat
## mu_0        4.97   0.01  0.51   4.32   4.98   5.63  9584   1
## mu_species[1] -0.15   0.01  0.51  -0.81  -0.15   0.50  9553   1
## mu_species[2] -0.05   0.01  0.51  -0.71  -0.06   0.60  9600   1
## mu_species[3]  0.40   0.01  0.51  -0.26   0.39   1.05  9565   1
## b_0         -0.06   0.01  0.50  -0.69  -0.07   0.59  9458   1
## b_species[1]   0.07   0.01  0.50  -0.58   0.08   0.70  9455   1
## b_species[2]   0.04   0.01  0.50  -0.61   0.04   0.67  9455   1
## b_species[3]  -0.17   0.01  0.50  -0.81  -0.16   0.46  9452   1
## sigma_mu_id   0.18   0.00  0.02   0.15   0.18   0.20  3189   1
## sigma_b_id    0.05   0.00  0.00   0.04   0.05   0.05  3830   1
## sigma_e[1]     0.98   0.00  0.01   0.97   0.98   0.99 28658   1
## sigma_e[2]     1.04   0.00  0.01   1.03   1.04   1.05 30303   1
## sigma_e[3]     0.94   0.00  0.01   0.93   0.94   0.95 22380   1
## lp__      5714.07  0.33 16.38 5692.69 5714.47 5734.66  2482   1
##

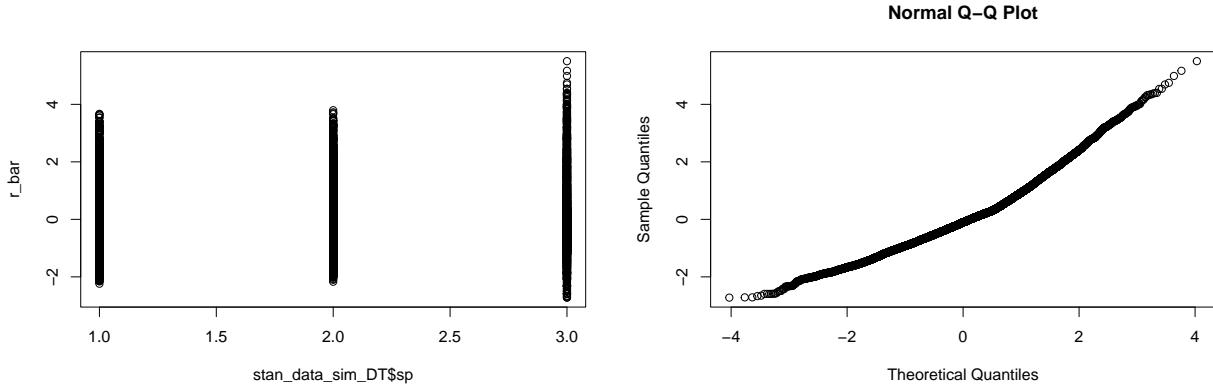
```

```

## Samples were drawn using NUTS(diag_e) at Fri Sep 17 10:23:14 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```





The first thing we note, based on the n_{eff} and Rhat values from the data summary and the traceplots of the parameter estimates across all four MCMC chains is that our model has converged across chains, and that chains are well mixed. We avoid interpreting the raw reports of parameter values in the data summary, as these don't capture parameter correlations, and most of our values of interest require combining estimated parameters, which we do further below.

The pairs plots indicate a great deal of correlation between the grand mean parameters for intercept and slope, and their corresponding species level adjustments. This is to be expected; as the estimated grand mean parameters move around the true value of species level parameters, estimates of the species level adjustments should change accordingly. Importantly, estimates for the grand mean of the intercept and the grand mean of the slope are not correlated. Additionally, there are not other irregularities in the parameter estimate distributions that would indicate the model was unidentifiable.

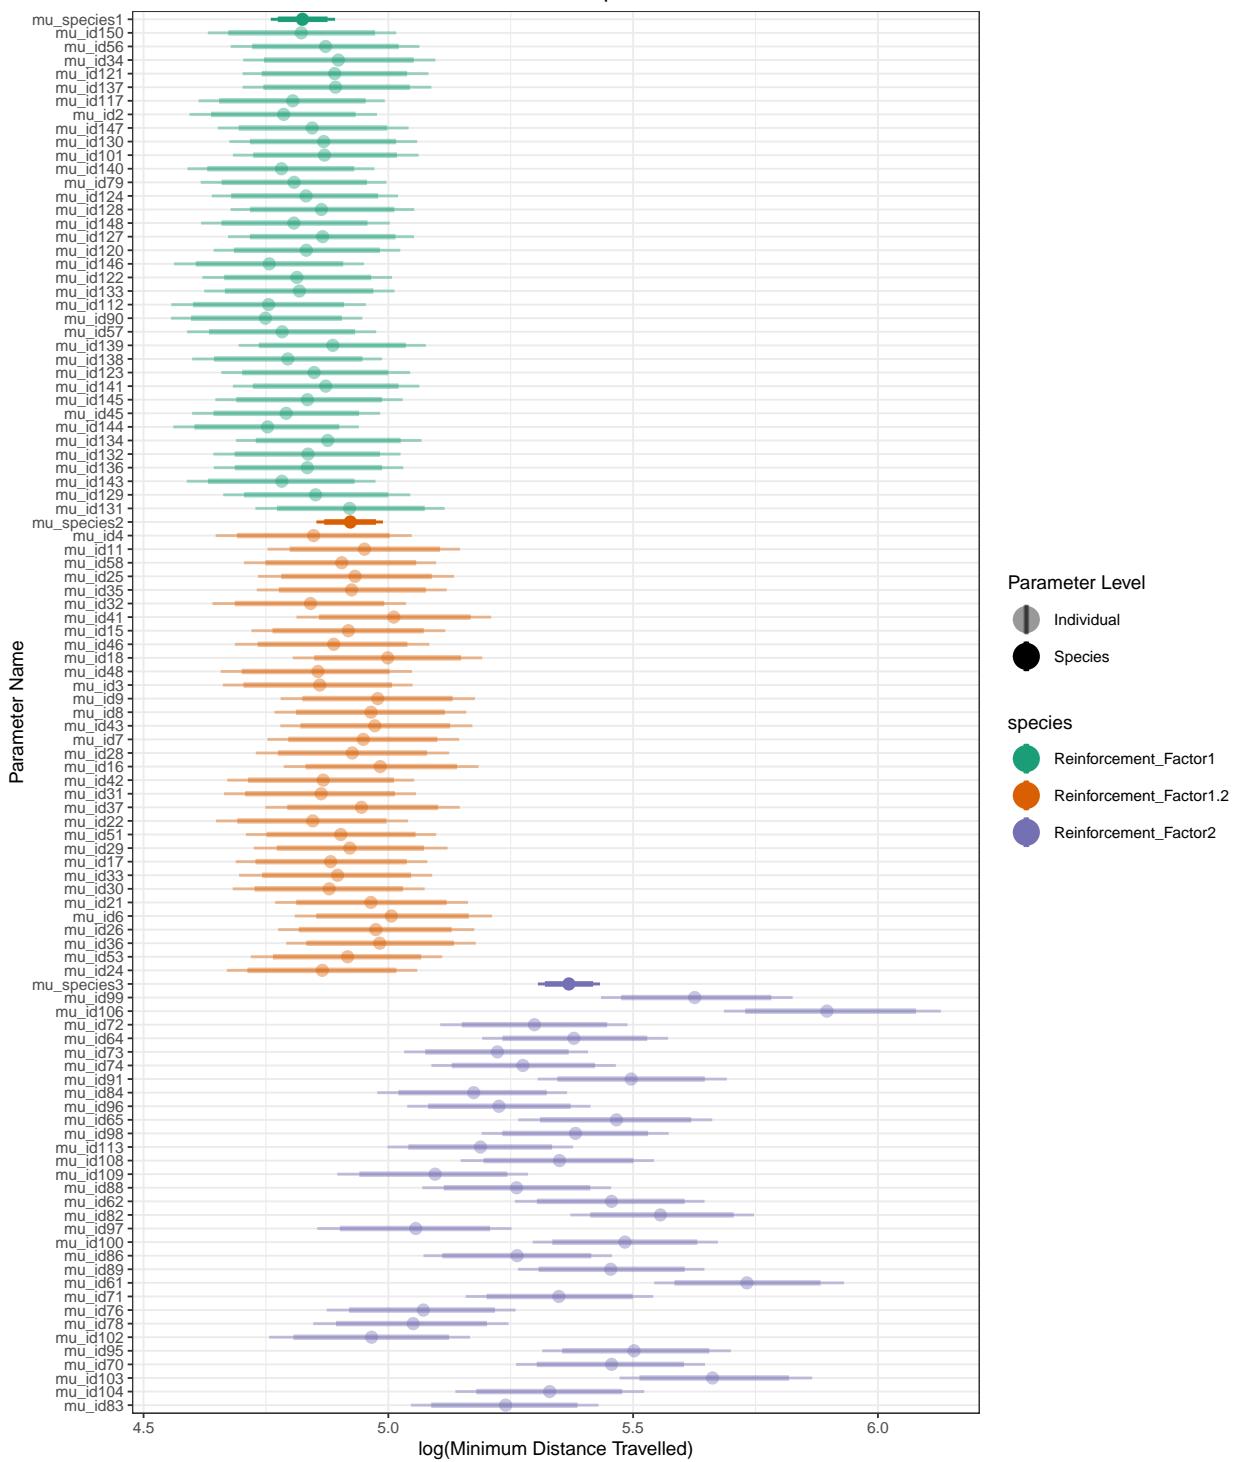
The residuals of this model, however, suggest our model does not fit the data well. Looking at residuals plotted with the predicted value on the x axis, there is a strong negative trend, indicated overestimation is more likely when predictions of distance traveled are low, and vice-versa. Additionally, the variance in residuals is much lower when predicted outcomes are low, reflecting that the convergence of Learning Factor 2 agents on short routes is not well modelled, and perhaps skewing parameter estimates. However, as our goal here is to understand these results in the context of our empirical analysis, we chose not to adjust this model, and instead consider the reasons why a power law function is a good model for actual primate learning in the Double Trapezoid array, but not for "learning" when the iterative reinforcement algorithm operates in the same array.

Thus, we move on by extracting the posteriors distributions of each parameter. First, we calculate species level intercepts and slopes by adding each species level adjustment to the grand means of the appropriate parameters. Then we do the same for individual-level parameter estimates, adding individual-level adjustments to the species level parameters just calculated.

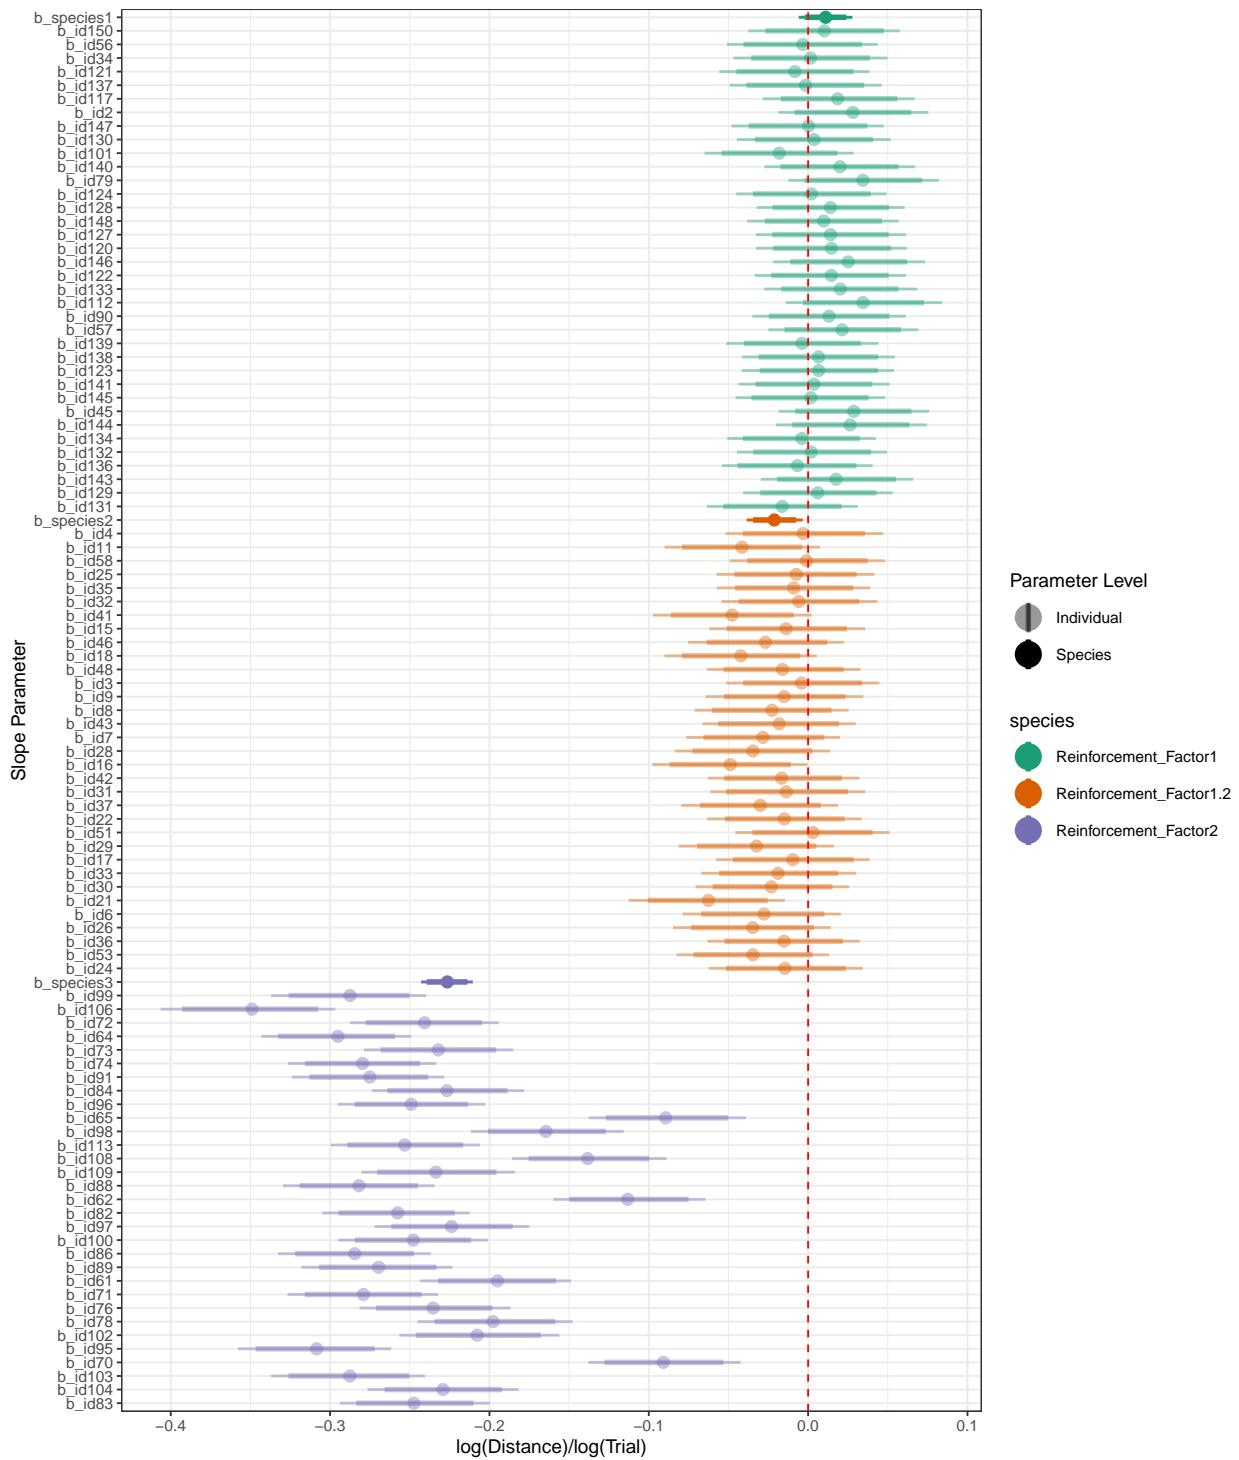
To visualize the resulting posteriors, we first determine quantiles of each posterior distribution and arrange these values in a data frame for easy plotting.

We then plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.

Estimated Distance Travelled on First Trial for Species and Individuals



Estimated Learning Rates for Species and Individuals



There are several interesting take-aways from these plots. First, there are significant differences in the estimated intercepts for agents of different reinforcement factors. Because we know the behavior of agents on the first trial to be identical, this must reflect the poor fit of the model. Considering the pattern of data, and the structure of the residuals, this is likely caused by a pattern of learning in which shortest route is learned very quickly once found, but the time it takes to find that route is quite stochastic. Real animals, on the other hand, show strategic exploration and gradual improvement - producing the power function learning

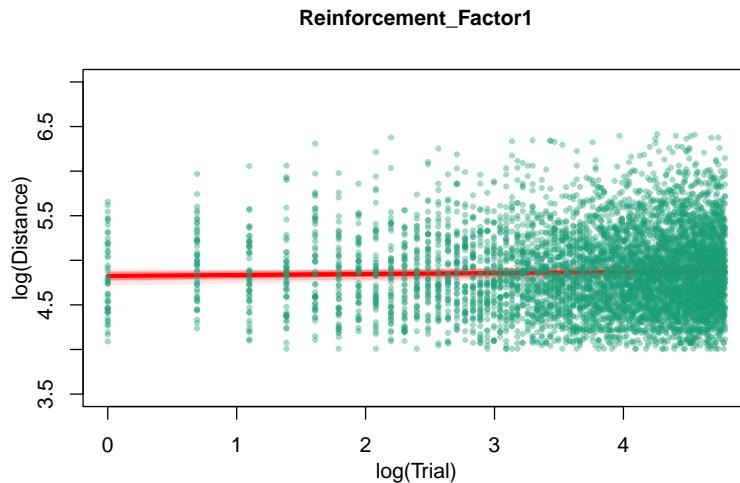
curves that fit the empirical data well. This difference in learning pattern for the simulated agents is also apparent in the high individual variation in estimated parameters for the learning factor 2. The learning rate is less a consequence of how fast the animal learns, but how long it takes to find a shorter route, and which route that is.

None-the-less, we do see overall learning rates that reflect our predictions about the iterative reinforcement algorithm. Though the estimated magnitude of the learning rate for reinforcement factor is small, more than 90% of the posterior estimates for the slope fall below zero, indicated improvement. For Reinforcement Factor 1, on the other hand, 0 is contained within the 80% Credible Interval, giving (as expected) no evidence of improvement. For Reinforcement Factor 2 agents, the magnitude of the slope parameter in the full posterior distribution is unsurprisingly much larger.

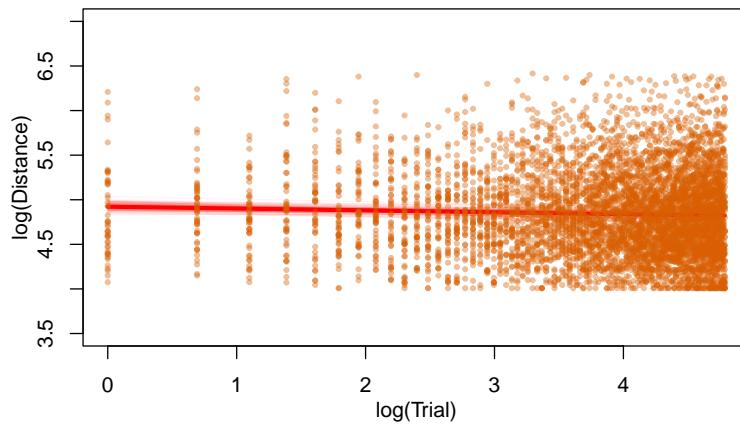
```
## pdf
## 2

## pdf
## 2
```

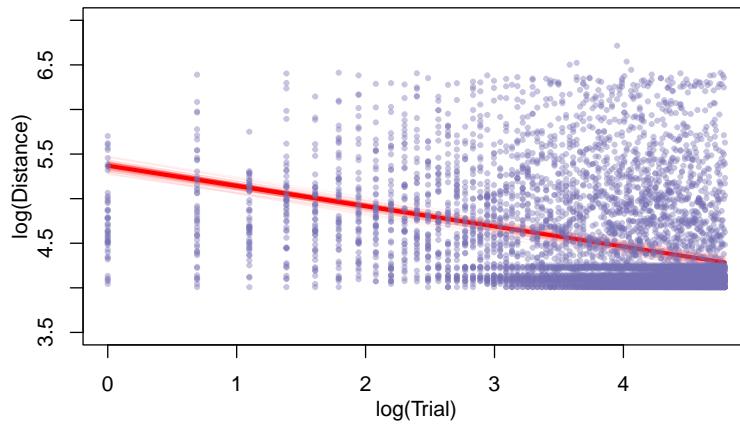
Finally, we visualize the mean regression models, with a sample of models from the posteriors and the raw data superimposed.



Reinforcement_Factor1.2

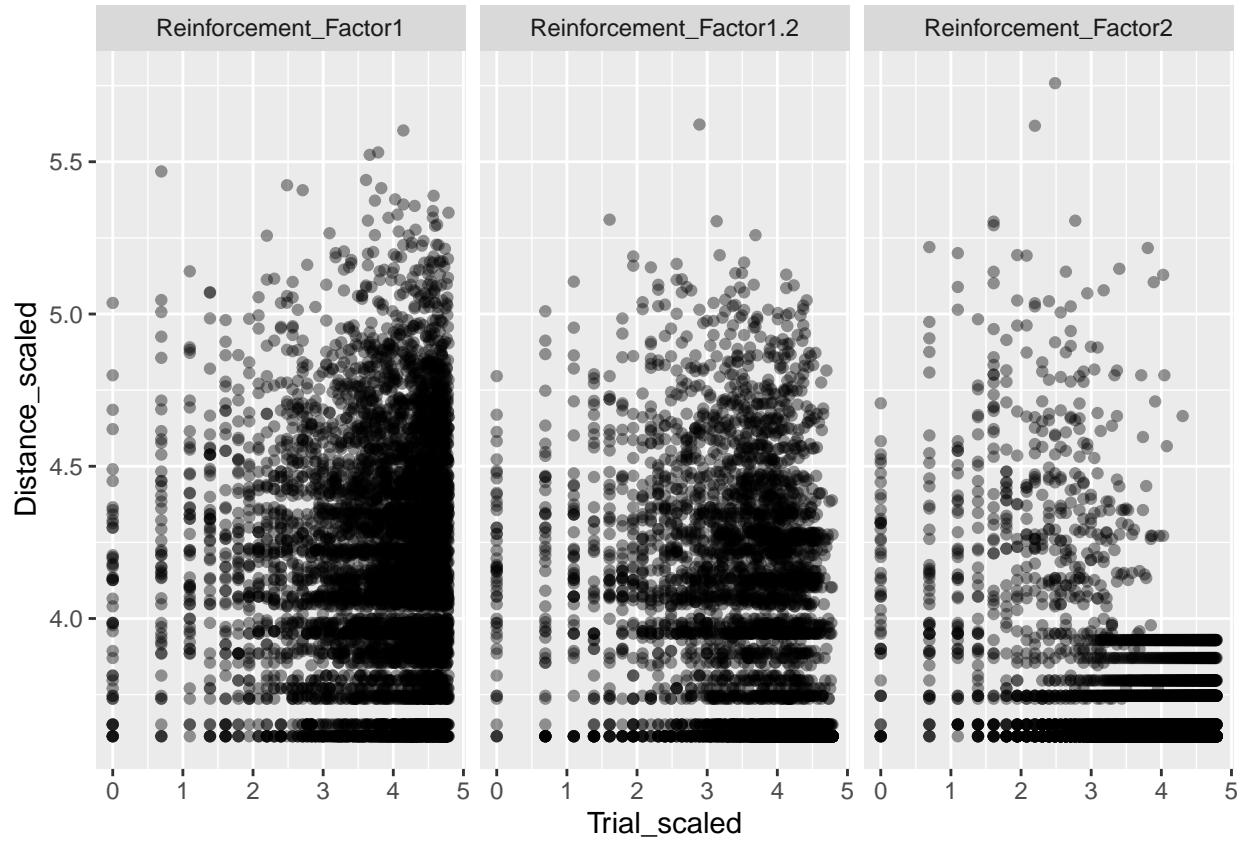


Reinforcement_Factor2



Simulation analysis, Pentagon

We set up the simulated data form the Pentagon array much as we did for the Double Trapezoid



```
## Warning: namespace 'raster' is not available and has been replaced
## by .GlobalEnv when processing object 'fit_sim'
```

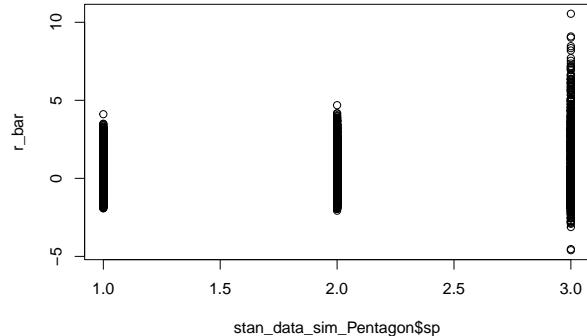
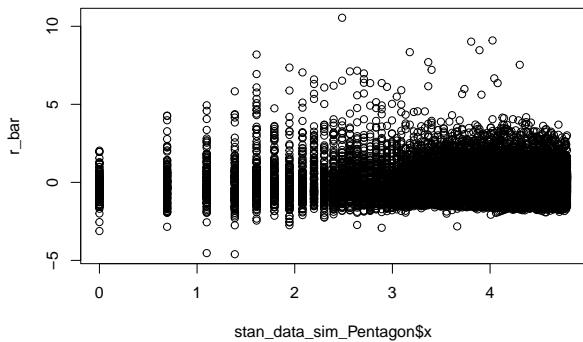
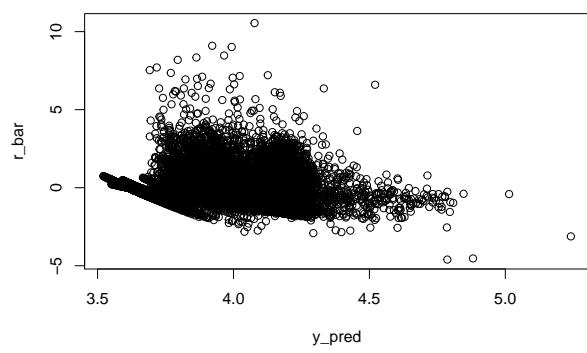
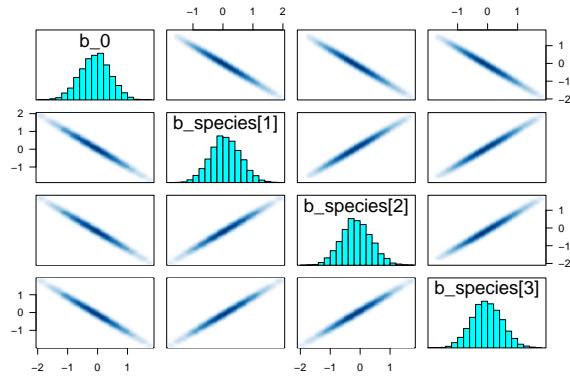
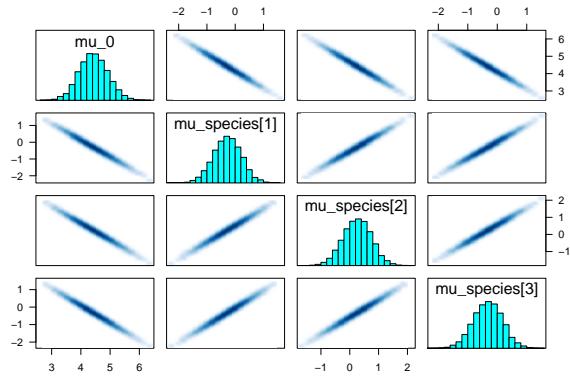
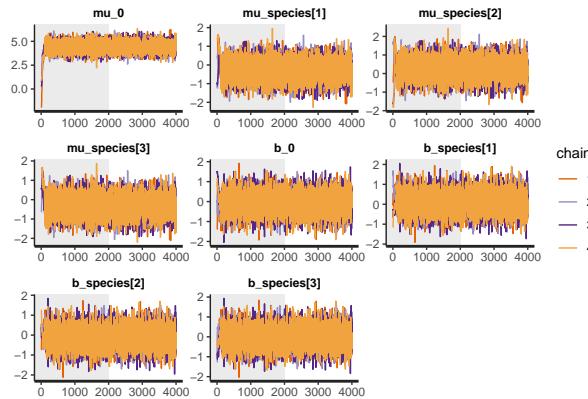
As with the Double Trapezoid, we check the model diagnostics after fitting it to the Pentagon simulations

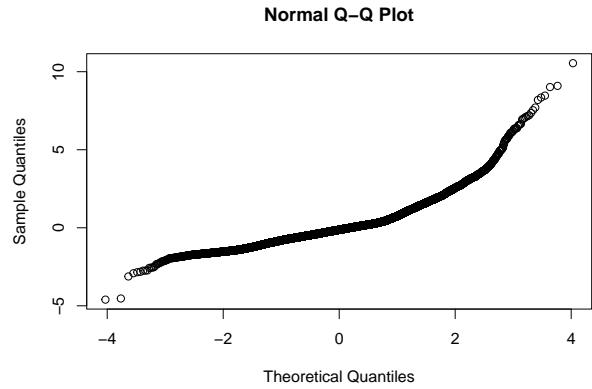
```
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=4000; warmup=2000; thin=1;
## post-warmup draws per chain=2000, total post-warmup draws=8000.
##
##           mean se_mean    sd    10%    50%    90% n_eff Rhat
## mu_0        4.43   0.01  0.49   3.82   4.42   5.06  3740   1
## mu_species[1] -0.28   0.01  0.49  -0.92  -0.28   0.33  3719   1
## mu_species[2]  0.26   0.01  0.49  -0.38   0.26   0.87  3765   1
## mu_species[3] -0.32   0.01  0.49  -0.95  -0.31   0.30  3719   1
## b_0         -0.08   0.01  0.50  -0.73  -0.07   0.56  3647   1
## b_species[1]  0.09   0.01  0.50  -0.55   0.08   0.73  3648   1
## b_species[2] -0.12   0.01  0.50  -0.77  -0.13   0.52  3641   1
## b_species[3] -0.02   0.01  0.50  -0.67  -0.03   0.62  3649   1
## sigma_mu_id  0.20   0.00  0.01   0.18   0.20   0.22  6191   1
## sigma_b_id   0.04   0.00  0.00   0.04   0.04   0.05  6052   1
## sigma_e[1]    0.83   0.00  0.01   0.82   0.83   0.84 18369   1
## sigma_e[2]    0.68   0.00  0.01   0.67   0.68   0.68 17986   1
## sigma_e[3]    0.30   0.00  0.00   0.29   0.30   0.30 17675   1
## lp__      16198.06  0.25 13.22 16181.34 16198.34 16214.79  2889   1
##
```

```

## Samples were drawn using NUTS(diag_e) at Tue Oct 05 10:07:41 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```



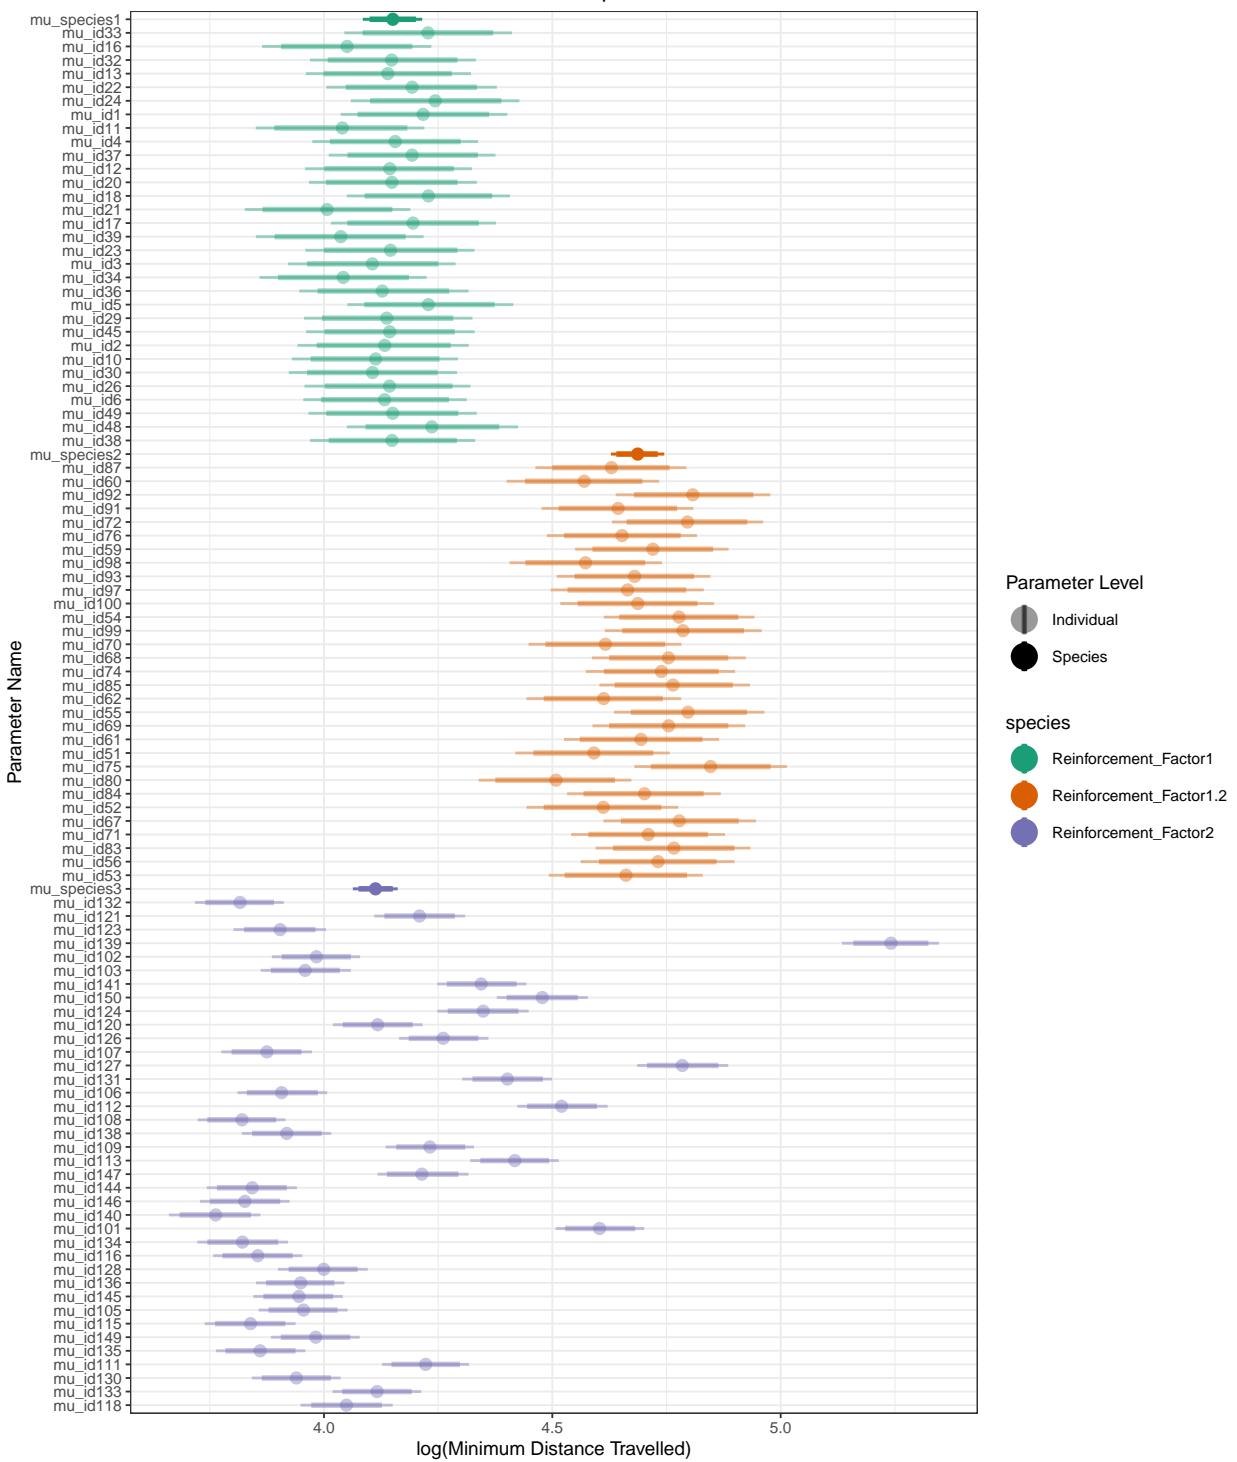


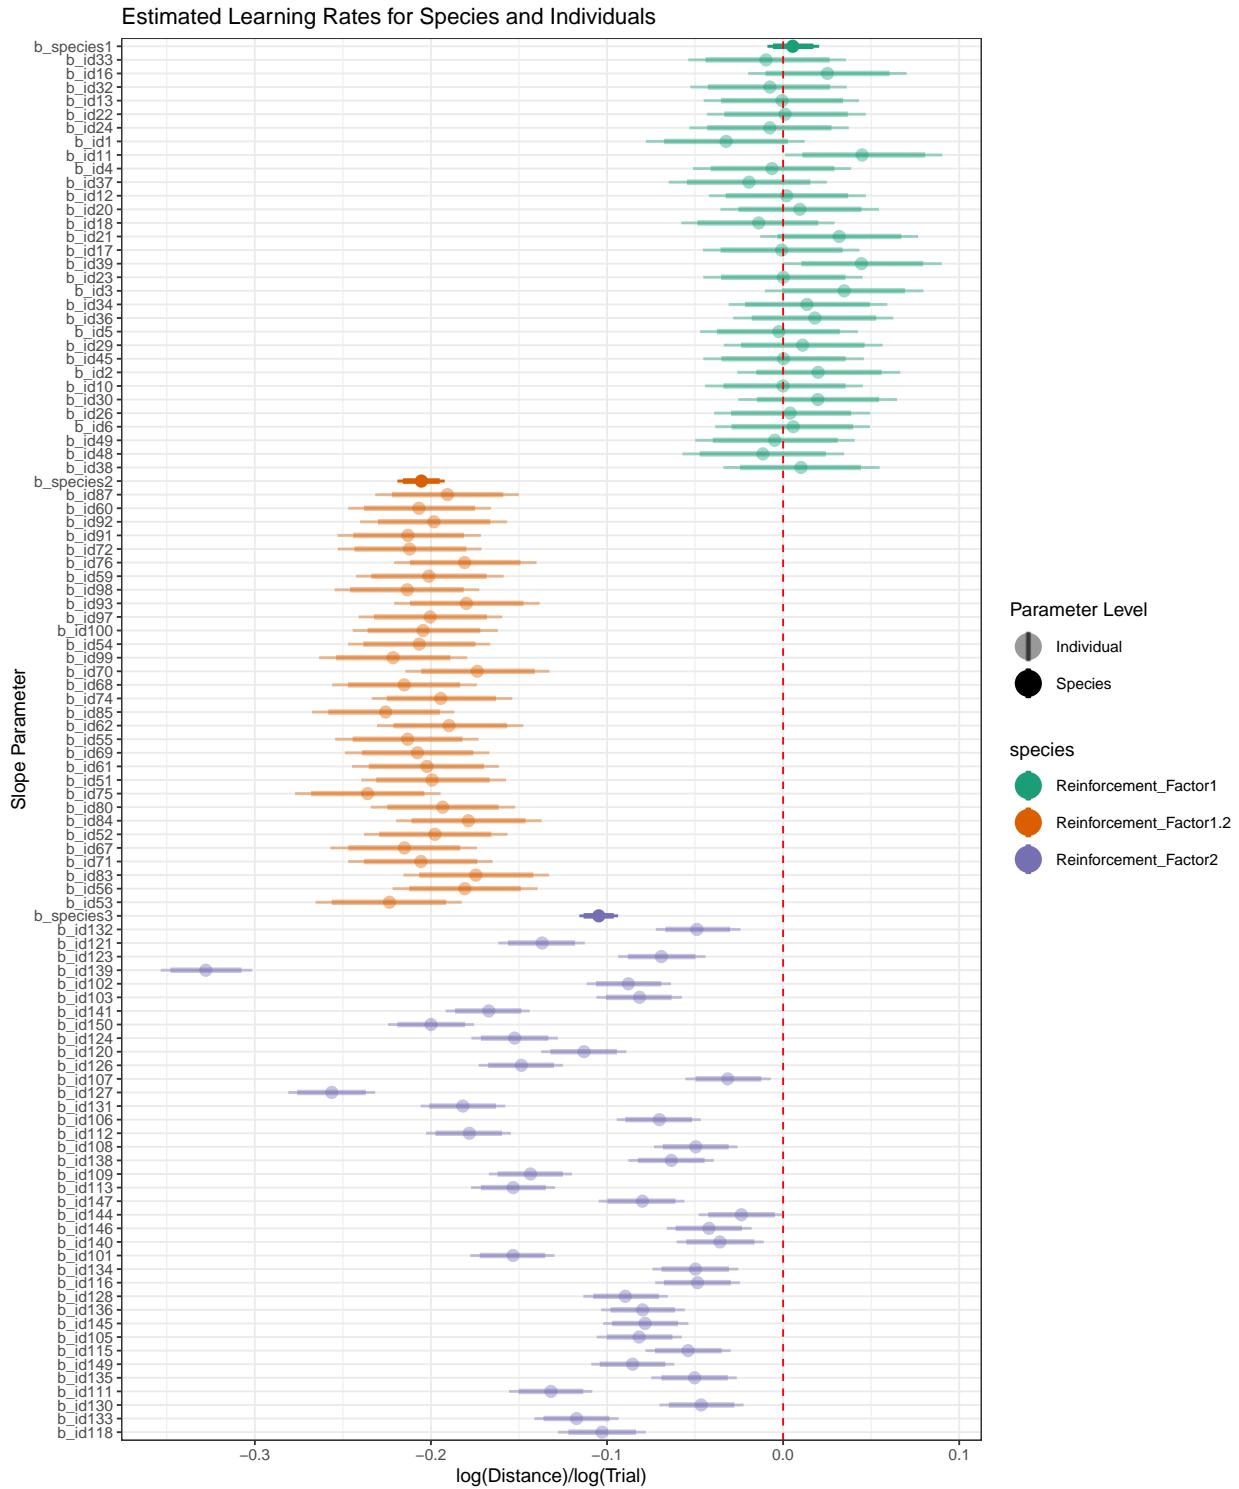
These results look much as they did for the double trapezoid: an identifiable, well-mixed model that fails to capture the real structure of the data, giving residuals that are biased by the Reinforcement Factor of the agent and the trial number.

Having assessed the fit of the model, we look at the parameter estimates, first getting species and individual level estimates by adding relevant deviations to mean estimates.

To visualize the resulting posteriors, we determine quantiles of each posterior distribution and plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.

Estimated Distance Travelled on First Trial for Species and Individuals

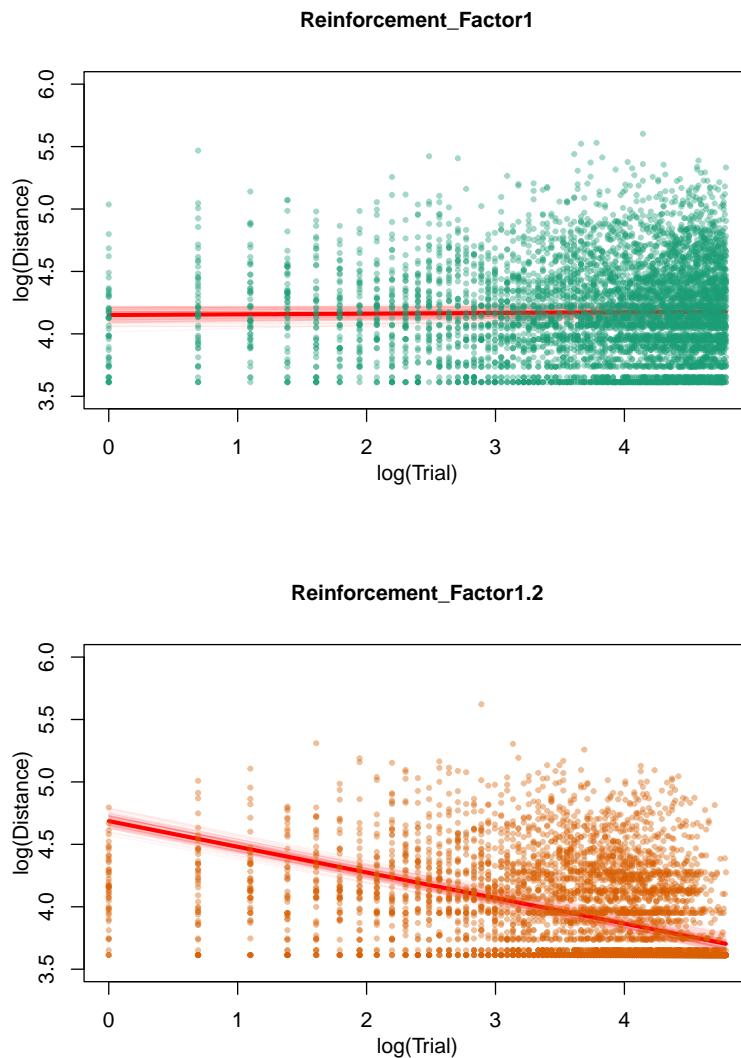




```
## pdf
## 2
```

```
## pdf
## 2
```

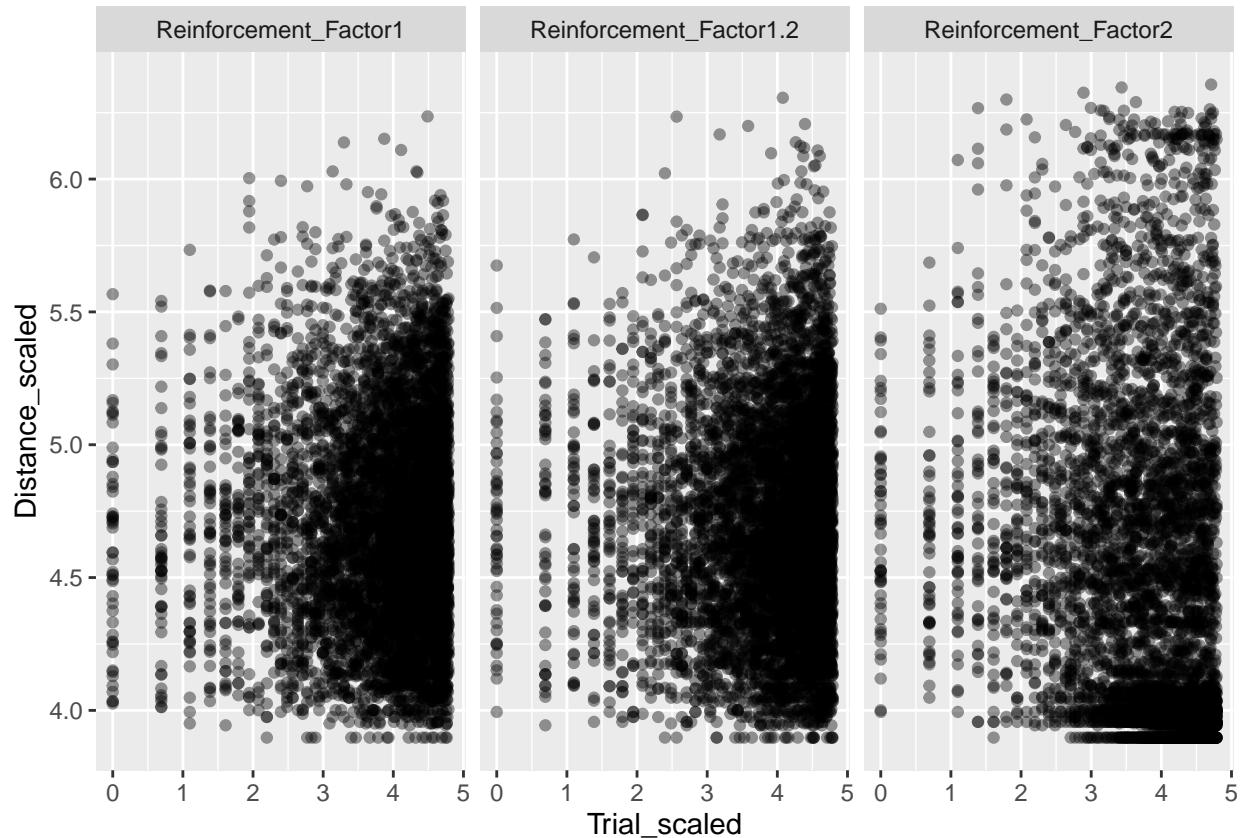
The estimated learning rate for reinforcement factor 1.2 is greater in magnitude than for reinforcement factor 2! Some data visualization is necessary to understand what is going on.





Simulation analysis, Z-Array

We set up the simulated data form the Pentagon array much as we did for the Double Trapezoid

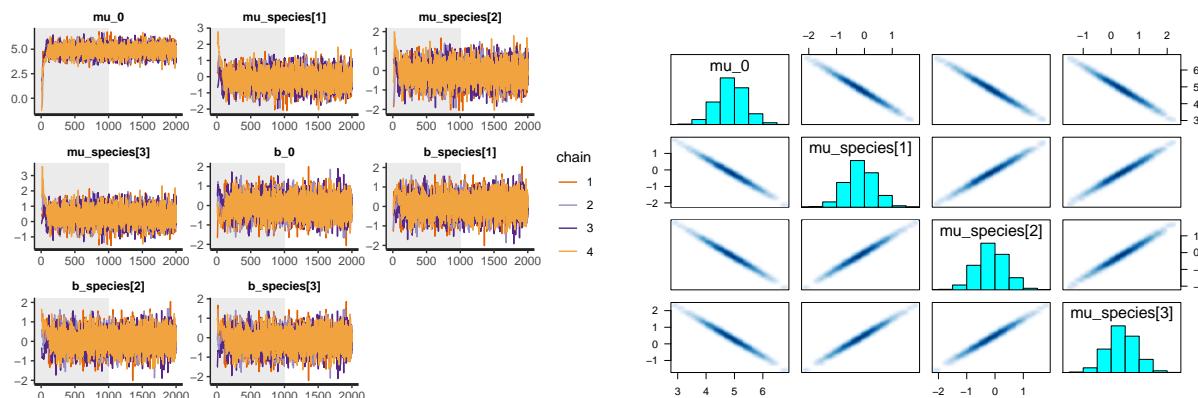


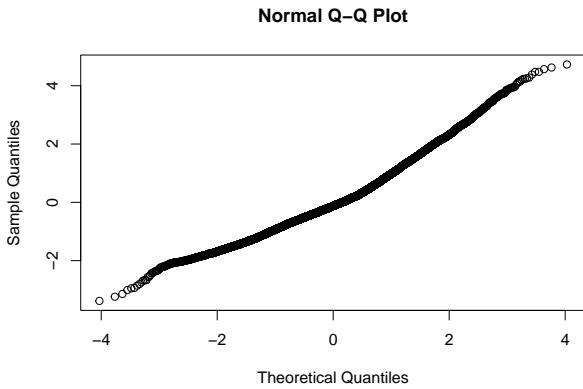
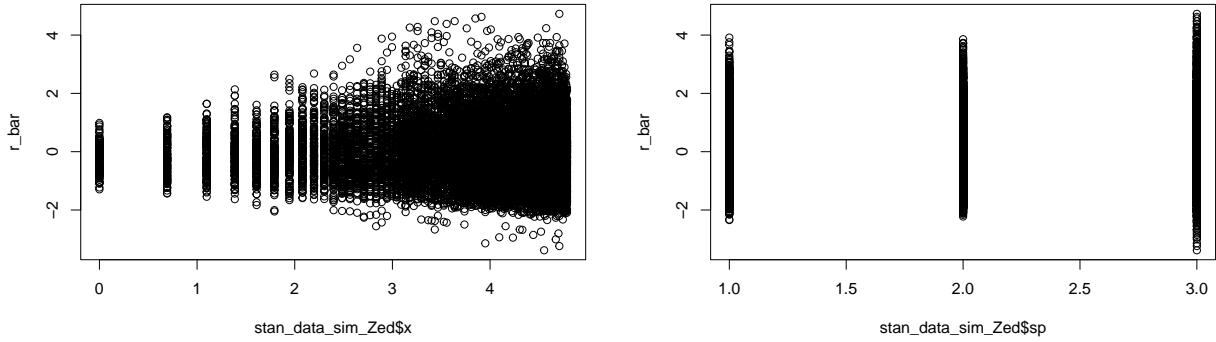
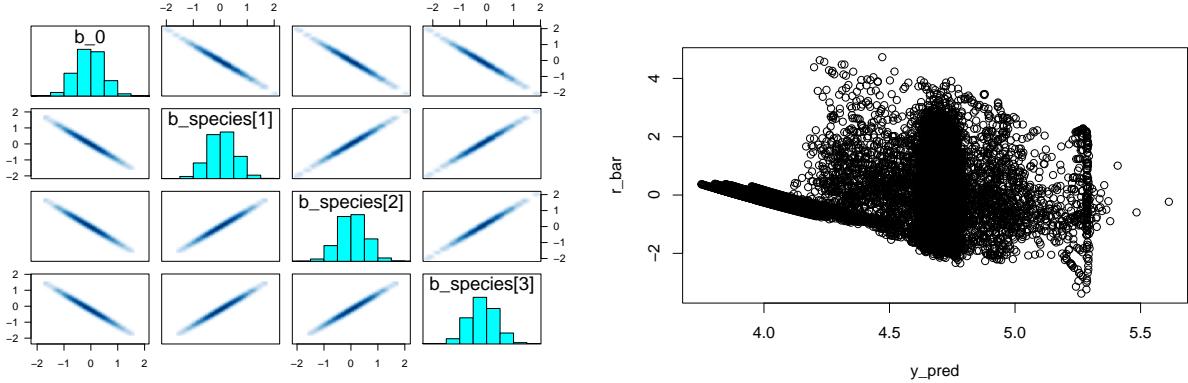
```
## Warning: namespace 'rgdal' is not available and has been replaced
## by .GlobalEnv when processing object 'fit_sim'
```

```
## Warning: namespace 'raster' is not available and has been replaced
## by .GlobalEnv when processing object 'fit_sim'
```

As with the Double Trapezoid, we check the model diagnostics after fitting it to the Zed simulations

```
## Inference for Stan model: DT_power_model.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##          mean se_mean    sd    10%    50%    90% n_eff Rhat
## mu_0      4.87    0.01  0.51   4.20   4.88   5.50 3028 1.00
## mu_species[1] -0.15    0.01  0.51  -0.79  -0.16   0.52 3018 1.00
## mu_species[2] -0.14    0.01  0.51  -0.76  -0.15   0.53 3025 1.00
## mu_species[3]  0.38    0.01  0.51  -0.26   0.38   1.05 3038 1.00
## b_0       -0.06    0.01  0.51  -0.70  -0.06   0.58 2660 1.00
## b_species[1]  0.05    0.01  0.51  -0.59   0.06   0.70 2665 1.00
## b_species[2]  0.05    0.01  0.51  -0.59   0.05   0.70 2658 1.00
## b_species[3] -0.15    0.01  0.51  -0.79  -0.16   0.49 2653 1.00
## sigma_mu_id  0.11    0.00  0.02   0.08   0.11   0.14 255  1.02
## sigma_b_id   0.05    0.00  0.00   0.05   0.05   0.06 937  1.01
## sigma_e[1]    0.89    0.00  0.01   0.88   0.89   0.90 9967 1.00
## sigma_e[2]    0.92    0.00  0.01   0.91   0.92   0.93 8928 1.00
## sigma_e[3]    0.95    0.00  0.01   0.94   0.95   0.96 4468 1.00
## lp__      6975.35  1.92 28.95 6942.57 6972.95 7010.55  228 1.03
##
## Samples were drawn using NUTS(diag_e) at Mon Oct 11 05:23:07 2021.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```



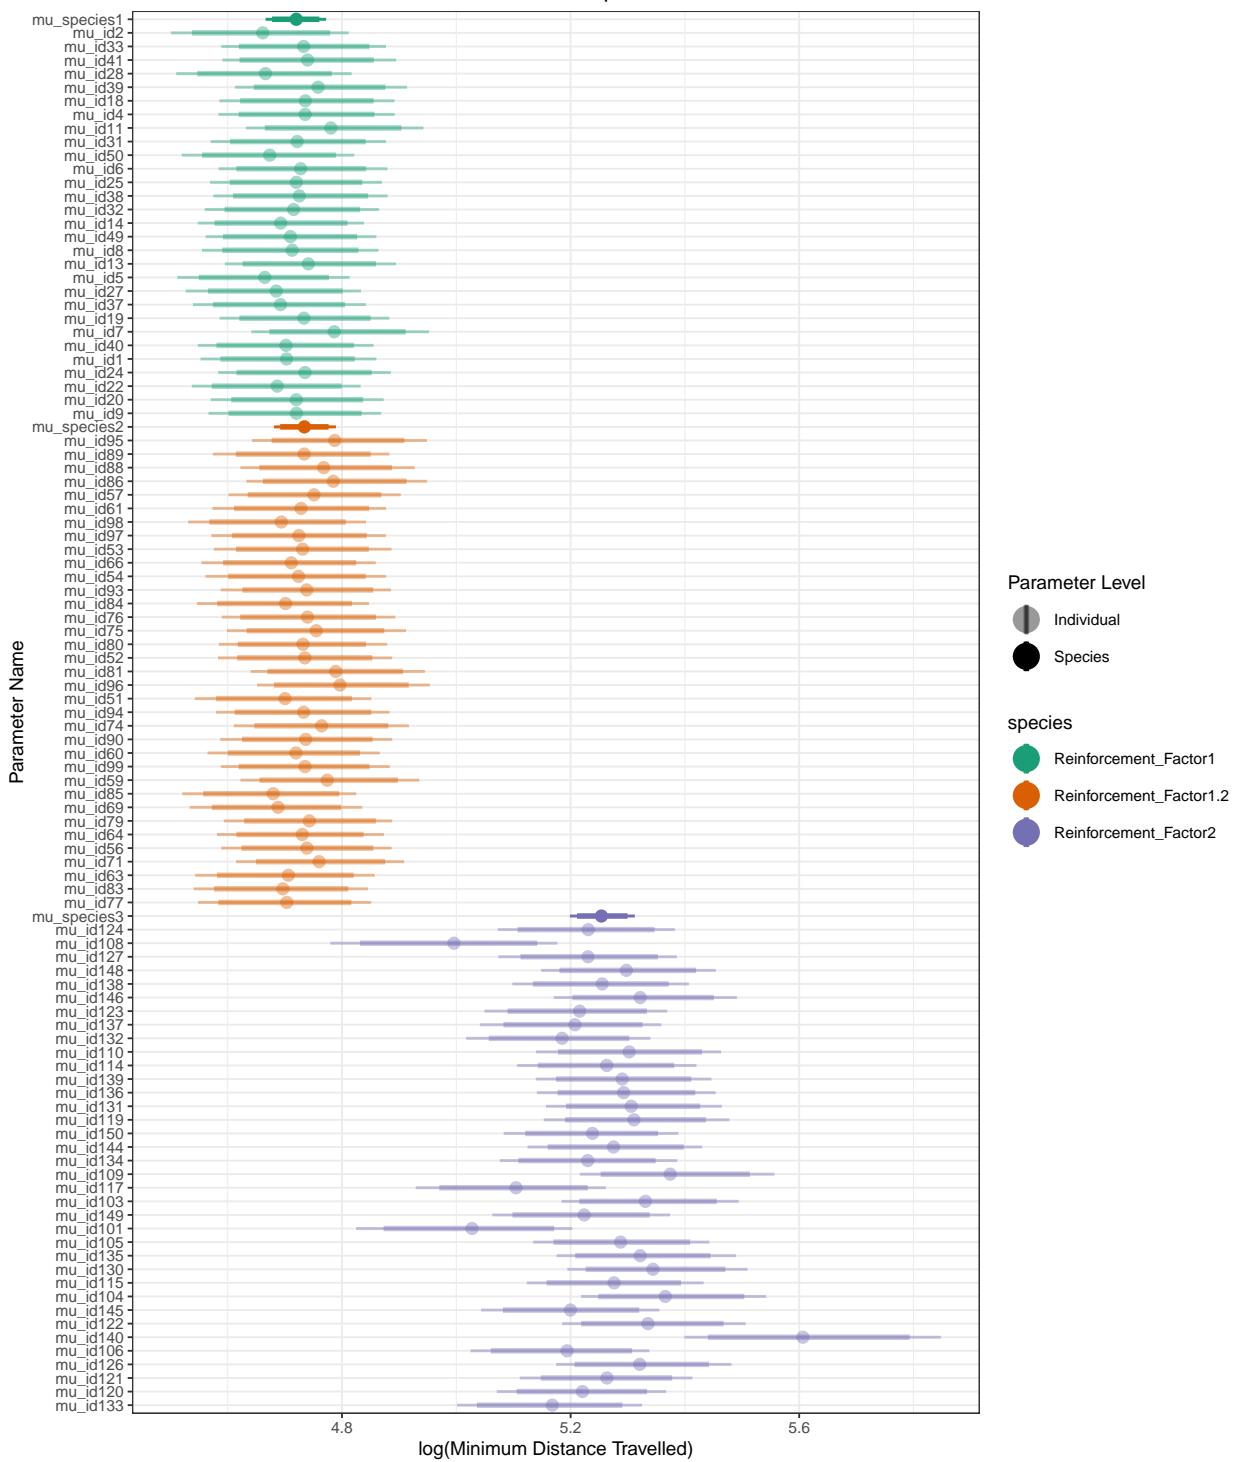


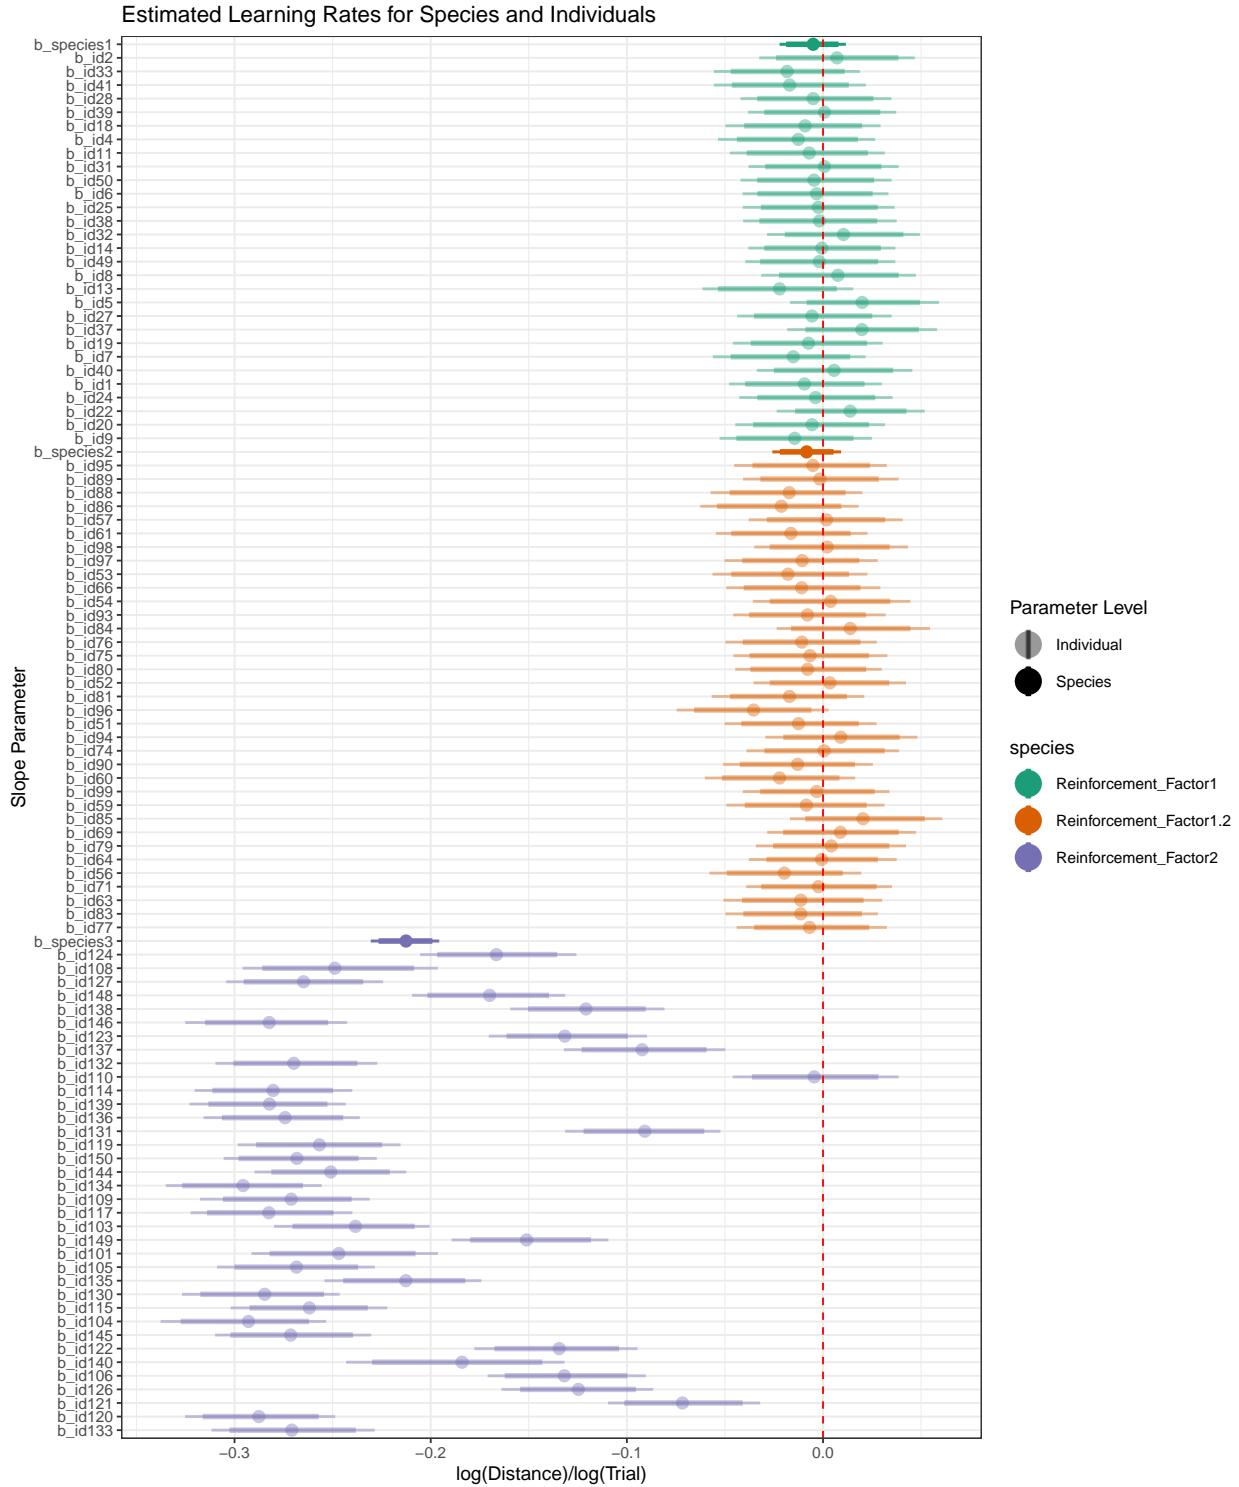
The MCMC chains for this model did not mix as well as in the DT or Pentagon simulations, however we deem the Rhat values and effective sample sizes to indicate good enough mixture and convergence to continue the analysis. As with the DT and Pentagon simulations, the learning model fails to capture some of the structure in the data, with residuals showing a downward trend along predicted distances travelled. Overall, though, residuals from the ZED array appear to be the most independently and identically distributed relative to the other simulation analyses.

Having assessed the fit of the model, we look at the parameter estimates, first getting species and individual level estimates by adding relevant deviations to mean estimates.

To visualize the resulting posteriors, we determine quantiles of each posterior distribution and plot the 80% (thick bars) and 90% (thin bars) confidence intervals for each species and individual level parameter.

Estimated Distance Travelled on First Trial for Species and Individuals

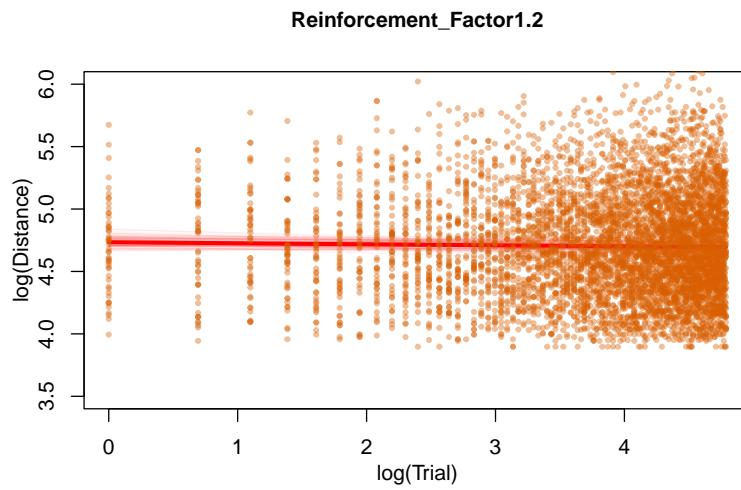
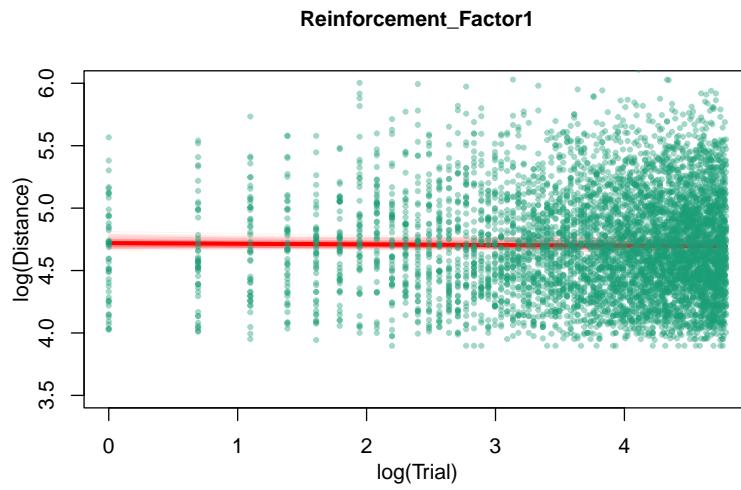


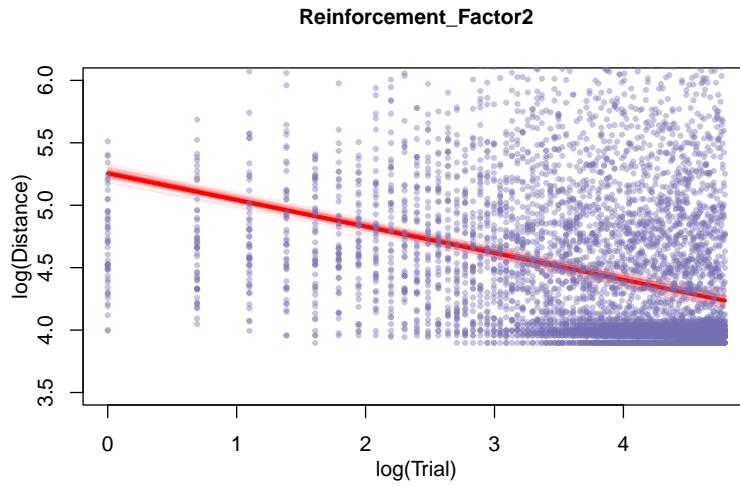


This results look very similar to the Double Trapezoid simulations. They have the same strong effect of reinforcement factor on learning rate, with an accompanying shift in estimated intercept indicating a mis-fit to the data. They also exhibit the high individual variance in parameter estimates for reinforcement factor 2 agents. Interestingly, a learning rate of 0 is contained in the 80% Credible Interval of reinforcement factor 1.2 agents, suggesting this level of reinforcement is not sufficient to learn shorter routes - a notable difference from the results in the Double Trapezoid.

```
## pdf  
## 2
```

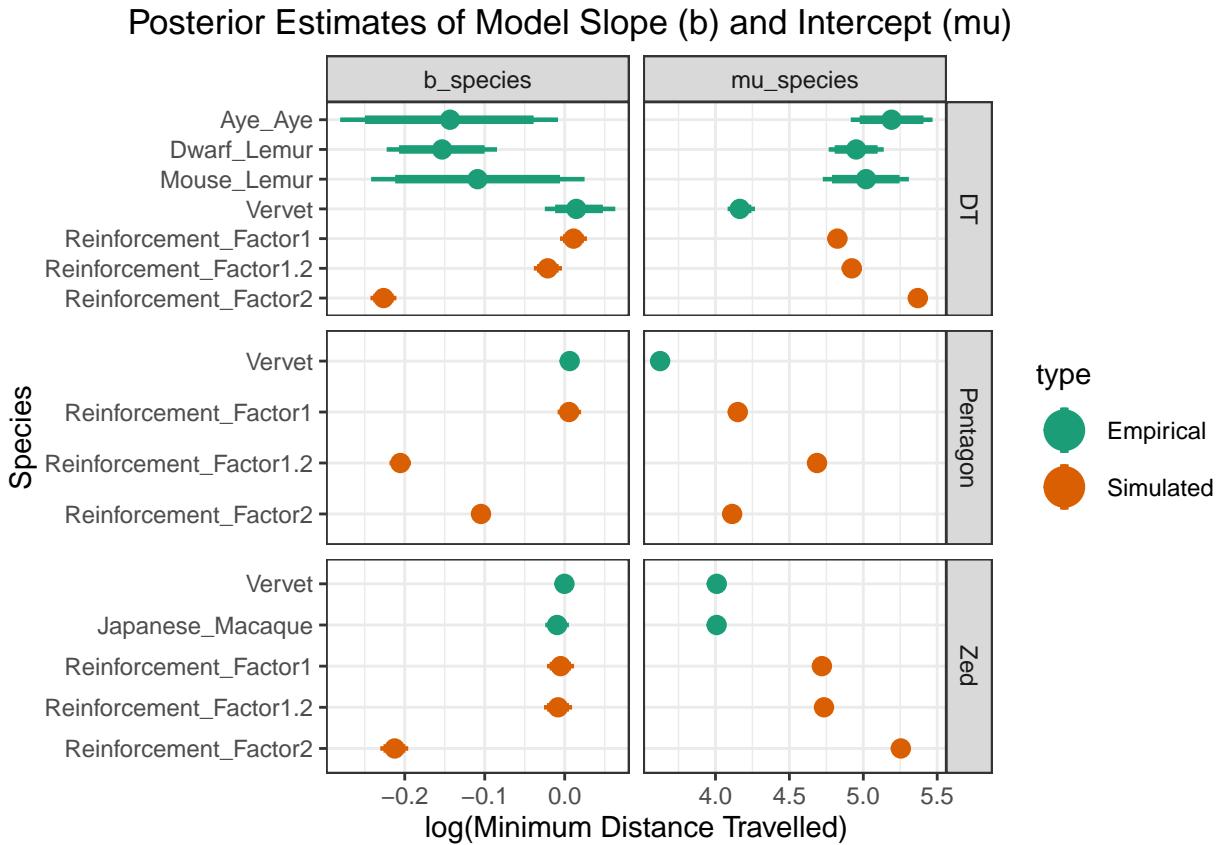
```
## pdf  
## 2
```





Integrated Analyses

First, we want to look at species level parameter estimates, for both real species and simulated agent types, within arrays.



```
## pdf
```

```
##    2
```

Second, we look at the posterior regressions for the empirical data (See Figure 2, manuscript)

```
## pdf  
## 2
```

And then the regression plots for the simulations (See Supp Figure 1, in DOCS)

```
## pdf  
## 2
```